


Bayes Classifiers

Adopted from slides by Alexander Ihler

Supervised Learning

- **Given** examples of a function $(X, Y = F(X))$
- **Find** function $\hat{Y} = h(X)$ to estimate $F(X)$
 - Continuous Y : Regression
 - Discrete Y : Classification

Nearest neighbor classifier



A basic classifier

- Training data $D = \{x^{(i)}, y^{(i)}\}$, Classifier $f(x)$
 - Discrete feature vector X
 - $f(x)$ is a contingency table
- Ex: credit rating prediction (bad/good)
 - X = income (low/med/high)
 - How can we make the most # of correct predictions?

| Features | # bad | # good |
|----------|-------|--------|
| X=0 | 42 | 15 |
| X=1 | 338 | 287 |
| X=2 | 3 | 5 |

A basic classifier

- Training data $D = \{x^{(i)}, y^{(i)}\}$, Classifier $f(x)$
 - Discrete feature vector X
 - $f(x)$ is a contingency table
- Ex: credit rating prediction (bad/good)
 - X = income (low/med/high)
 - How can we make the most # of correct predictions?
- Predict more likely outcome
for each possible observation

| Features | # bad | # good |
|----------|-------|--------|
| X=0 | 42 | 15 |
| X=1 | 338 | 287 |
| X=2 | 3 | 5 |

A basic classifier


- Training data $D = \{x^{(i)}, y^{(i)}\}$, Classifier $f(x)$
 - Discrete feature vector X
 - $f(x)$ is a contingency table
- Ex: credit rating prediction (bad/good)
 - X = income (low/med/high)
 - How can we make the most # of correct predictions?
- Predict more likely outcome for each possible observation
- Can normalize into probability:
 $p(y = \text{good} \mid X = x)$
- How to generalize?

| Features | # bad | # good |
|----------|-------|--------|
| X=0 | 42 | 15 |
| X=1 | 338 | 287 |
| X=2 | 3 | 5 |

Bayes rule

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

- Two events: headache, flu
- $p(H) = 1/10$
- $p(F) = 1/40$
- $p(H|F) = 1/2$




- You wake up with a headache – what is the chance that you have the flu?

Bayes rule

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$


- Two events: headache, flu
- $p(H) = 1/10$
- $p(F) = 1/40$
- $p(H|F) = 1/2$
- $P(H \& F) = ?$
- $P(F|H) = ?$



Bayes rule

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$


- Two events: headache, flu
- $p(H) = 1/10$
- $p(F) = 1/40$
- $p(H|F) = 1/2$
- $P(H \& F) = p(F) p(H|F)$
 $= (1/2) * (1/40) = 1/80$
- $P(F|H) = ?$



Bayes rule

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

- Two events: headache, flu
- $p(H) = 1/10$
- $p(F) = 1/40$
- $p(H|F) = 1/2$
- $P(H \& F) = p(F) p(H|F)$
 $= (1/2) * (1/40) = 1/80$
- $P(F|H) = p(H \& F) / p(H)$
 $= (1/80) / (1/10) = 1/8$



Classification and probability

- Suppose we want to model the data
- Prior probability of each class, $p(y)$
 - E.g., fraction of applicants that have good credit
- Distribution of features given the class, $p(x | y = c)$
 - How likely are we to see “ x ” in users with good credit?

- Joint distribution
$$p(y|x)p(x) = p(x, y) = p(x|y)p(y)$$

- Bayes Rule:
$$\Rightarrow p(y|x) = p(x|y)p(y)/p(x)$$

(Use the rule of total probability to calculate the denominator!) \longrightarrow

$$= \frac{p(x|y)p(y)}{\sum_c p(x|y = c)p(y = c)}$$

Bayes classifiers


- Training data
 - Estimate $p(y = c)$
 - Split by class
 - $D_c = \{x^{(i)}: y^{(i)} = c\}$
- Estimate $p(x | y = c)$ using D_c
- Estimate $p(y | x)$ using Bayes rule
- For a discrete X , this recalculates the same table...

| Features | # bad | # good | | | p(x y=0) | p(x y=1) | | | p(y=0 x) | p(y=1 x) |
|----------|-------|--------|--|--|------------|------------|--|--|----------|----------|
| X=0 | 42 | 15 | | | 42 / 383 | 15 / 307 | | | .7368 | .2632 |
| X=1 | 338 | 287 | | | 338 / 383 | 287 / 307 | | | .5408 | .4592 |
| X=2 | 3 | 5 | | | 3 / 383 | 5 / 307 | | | .3750 | .6250 |

| | | |
|------|---------|---------|
| p(y) | 383/690 | 307/690 |
|------|---------|---------|

Bayes classifiers


- Training data
 - Estimate $p(y = c)$
 - Split by class
 - $D_c = \{x^{(i)}: y^{(i)} = c\}$
- Estimate $p(x | y = c)$ using D_c
- Estimate $p(y | x)$ using Bayes rule
- For continuous X , can use any density estimate like
 - Histogram
 - Gaussian
 - ...



Gaussian models

- Estimate parameters of the Gaussians from the data



$$\alpha = \frac{m_1}{m} = \hat{p}(y = c_1) \quad \hat{\mu} = \frac{1}{m} \sum_j x^{(j)} \quad \hat{\sigma}^2 = \frac{1}{m} \sum_j (x^{(j)} - \mu)^2$$



Multivariate Gaussian models

- Similar to univariate case

$$\mathcal{N}(\underline{x} ; \underline{\mu}, \Sigma) = \frac{1}{(2\pi)^{d/2}} |\Sigma|^{-1/2} \exp \left\{ -\frac{1}{2} (\underline{x} - \underline{\mu})^T \Sigma^{-1} (\underline{x} - \underline{\mu}) \right\}$$



Example: Gaussian Bayes for Iris Data


- Fit Gaussian distribution to each class $\{0,1,2\}$

$$p(y) = \text{Discrete}(\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$$

$$p(x_1, x_2 | y = 0) = \mathcal{N}(x; \mu_0, \Sigma_0)$$

$$p(x_1, x_2 | y = 1) = \mathcal{N}(x; \mu_1, \Sigma_1)$$

$$p(x_1, x_2 | y = 2) = \mathcal{N}(x; \mu_2, \Sigma_2)$$



Bayes Classifiers: Naïve Bayes

Bayes classifiers

- Estimate $p(y) = [p(y = 0), p(y = 1) \dots]$
- Estimate $p(x | y = c)$ for each class c
- Calculate $p(y = c | x)$ using Bayes rule
- Choose the most likely class c
- For a discrete X , can represent as a contingency table...
 - What about if we have more discrete features?

| Features | # bad | # good | | $p(x \mid y=0)$ | $p(x \mid y=1)$ | | $p(y=0 \mid x)$ | $p(y=1 \mid x)$ |
|----------|---------|---------|---------------|-----------------|-----------------|---------------|-----------------|-----------------|
| X=0 | 42 | 15 | \Rightarrow | 42 / 383 | 15 / 307 | \Rightarrow | .7368 | .2632 |
| X=1 | 338 | 287 | | 338 / 383 | 287 / 307 | | .5408 | .4592 |
| X=2 | 3 | 5 | | 3 / 383 | 5 / 307 | | .3750 | .6250 |
| $p(y)$ | 383/690 | 307/690 | | | | | | |

Joint distributions

- Make a truth table of all combinations of values

| A | B | C |
|---|---|---|
| 0 | 0 | 0 |
| 0 | 0 | 1 |
| 0 | 1 | 0 |
| 0 | 1 | 1 |
| 1 | 0 | 0 |
| 1 | 0 | 1 |
| 1 | 1 | 0 |
| 1 | 1 | 1 |

Joint distributions

- Make a truth table of all combinations of values
- For each combination of values, determine how probable it is
- Total probability must sum to one
- How many values did we specify?

| A | B | C | $p(A,B,C \mid y=1)$ |
|---|---|---|---------------------|
| 0 | 0 | 0 | 0.4 |
| 0 | 0 | 1 | 0.1 |
| 0 | 1 | 0 | 0.0 |
| 0 | 1 | 1 | 0.0 |
| 1 | 0 | 0 | 0.1 |
| 1 | 0 | 1 | 0.2 |
| 1 | 1 | 0 | 0.1 |
| 1 | 1 | 1 | 0.1 |

Overfitting & density estimation

- Estimate probabilities from the data
 - E.g., how many times (what fraction) did each outcome occur?
- M data $\ll 2^N$ parameters?
- What about the zeros?
 - We learn that certain combinations are impossible?
 - What if we see these later in test data?
- Overfitting!

| A | B | C | $p(A,B,C \mid y=1)$ |
|---|---|---|---------------------|
| 0 | 0 | 0 | 4/10 |
| 0 | 0 | 1 | 1/10 |
| 0 | 1 | 0 | 0/10 |
| 0 | 1 | 1 | 0/10 |
| 1 | 0 | 0 | 1/10 |
| 1 | 0 | 1 | 2/10 |
| 1 | 1 | 0 | 1/10 |
| 1 | 1 | 1 | 1/10 |


Overfitting & density estimation

- Reduce the model complexity
 - E.g., assume that features are conditionally independent of one another given the class label
- Conditional Independence:
- $p(a,b,c|y) = p(a|y) p(b|y) p(c|y)$
- $p(x_1, x_2, \dots x_N | y=1) = p(x_1 | y=1) p(x_2 | y=1) \dots p(x_N | y=1)$
- Only need to estimate each individually

| A | p(A y=1) |
|---|------------|
| 0 | .4 |
| 1 | .6 |

| B | p(B y=1) |
|---|------------|
| 0 | .7 |
| 1 | .3 |

| C | p(C y=1) |
|---|------------|
| 0 | .1 |
| 1 | .9 |



| A | B | C | p(A,B,C y=1) |
|---|---|---|----------------|
| 0 | 0 | 0 | .4 * .7 * .1 |
| 0 | 0 | 1 | .4 * .7 * .9 |
| 0 | 1 | 0 | .4 * .3 * .1 |
| 0 | 1 | 1 | ... |
| 1 | 0 | 0 | |
| 1 | 0 | 1 | |
| 1 | 1 | 0 | |
| 1 | 1 | 1 | |

Naïve Bayes Models

- Naïve Bayes:
 - $p(y | \mathbf{x}) = p(\mathbf{x} | y) p(y) / p(\mathbf{x})$
 - Estimate $p(y)$ for each class y
 - $p(\mathbf{x} | y) = \prod_i p(x_i | y)$
 - Estimate $p(x_i | y)$ for each feature x_i and class y

Predict $y = c_1$ if $p(\mathbf{x} | y = c_1) p(y = c_1) > p(\mathbf{x} | y = c_2) p(y = c_2)$

- Note: may not be a good model of the data
 - Doesn't capture correlations in \mathbf{x} 's
 - Can't capture some dependencies
- But in practice it often does quite well!

Example: Naïve Bayes

Observed Data:

| x_1 | x_2 | y |
|-------|-------|-----|
| 1 | 1 | 0 |
| 1 | 0 | 0 |
| 1 | 0 | 1 |
| 0 | 0 | 0 |
| 0 | 1 | 1 |
| 1 | 1 | 0 |
| 0 | 0 | 1 |
| 1 | 0 | 1 |

$$\hat{p}(y = 1) = \frac{4}{8} = (1 - \hat{p}(y = 0))$$

$$\hat{p}(x_1 = 1|y = 0) = \frac{3}{4}$$

$$\hat{p}(x_2 = 1|y = 0) = \frac{2}{4}$$

$$\hat{p}(x_1 = 1|y = 1) = \frac{2}{4}$$

$$\hat{p}(x_2 = 1|y = 1) = \frac{1}{4}$$

$$\begin{aligned}\hat{p}(x_1 = 1, x_2 = 1|y = 1) &= \hat{p}(x_1 = 1|y = 1) \hat{p}(x_2 = 1|y = 1) \\ &= \frac{2}{4} \times \frac{1}{4}\end{aligned}$$

Example: Naïve Bayes

Observed Data:

| x_1 | x_2 | y |
|-------|-------|-----|
| 1 | 1 | 0 |
| 1 | 0 | 0 |
| 1 | 0 | 1 |
| 0 | 0 | 0 |
| 0 | 1 | 1 |
| 1 | 1 | 0 |
| 0 | 0 | 1 |
| 1 | 0 | 1 |

$$\hat{p}(y = 1) = \frac{4}{8} = (1 - \hat{p}(y = 0))$$

$$\hat{p}(x_1 = 1|y = 0) = \frac{3}{4}$$

$$\hat{p}(x_2 = 1|y = 0) = \frac{2}{4}$$

$$\hat{p}(x_1 = 1|y = 1) = \frac{2}{4}$$

$$\hat{p}(x_2 = 1|y = 1) = \frac{1}{4}$$

Prediction given some observation x ?

$$\hat{p}(y = 1)\hat{p}(x = 11|y = 1) \begin{matrix} < \\ > \end{matrix} \hat{p}(y = 0)\hat{p}(x = 11|y = 0)$$

$$\frac{4}{8} \times \frac{2}{4} \times \frac{1}{4} \qquad \frac{4}{8} \times \frac{3}{4} \times \frac{2}{4}$$

Decide class 0

Naïve Bayes Models for Spam

- $y \in \{spam, not\ spam\}$
- X = observed words in email
 - Ex: ["the" ... "probabilistic" ... "lottery" ...]
 - "1" if word appears; "0" if not
- 1000's of possible words: 2^{1000s} parameters?
- # of atoms in the universe: » 2^{270} ...
- Model words **given** email type as independent
- Some words more likely for spam ("lottery")
- Some more likely for real ("probabilistic")
- Only 1000's of parameters now...

Summary

- Bayes rule; $p(y | x)$
- Bayes classifiers
 - Learn $p(x | y = C)$, $p(y = C) \Rightarrow p(y = C | x)$
- Naïve Bayes classifiers
 - Assume features are independent given class:
$$p(\boldsymbol{x} | y = C) = p(x_1 | y = C) p(x_2 | y = C) \dots$$
- Maximum likelihood (empirical) estimators for
 - Discrete variables
 - Gaussian variables
 - Overfitting