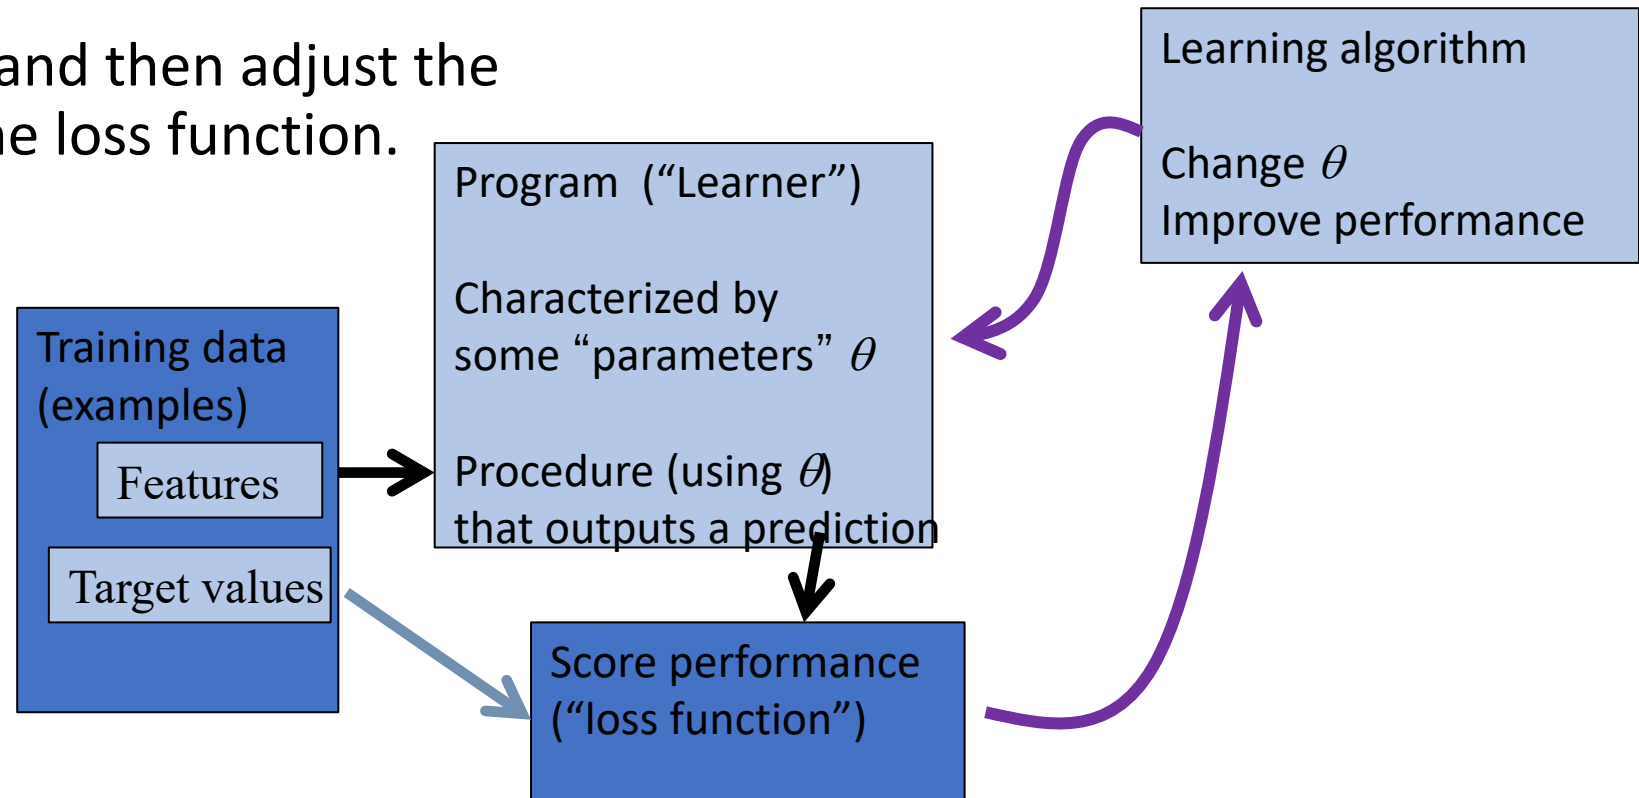


Linear Classifiers

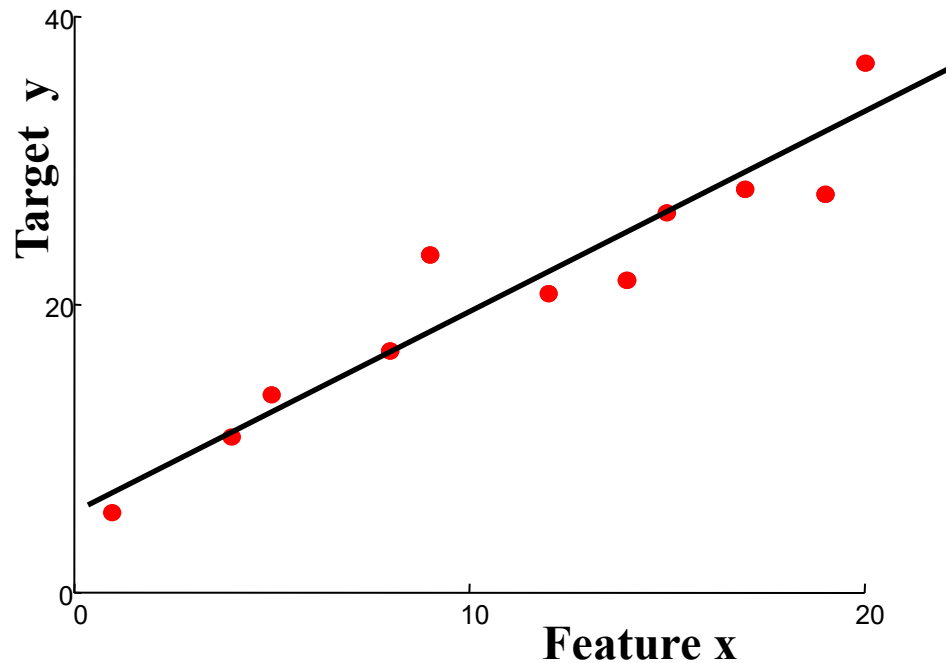
Adopted from slides by Alexander Ihler

Supervised Learning

- **Given** examples of a function ($X, Y = F(X)$)
- **Find** function $\hat{Y} = h(X)$ to estimate $F(X)$
 - Discrete Y : Classification
- Formulate a loss function and then adjust the parameters to minimize the loss function.



Linear regression



“Predictor”:

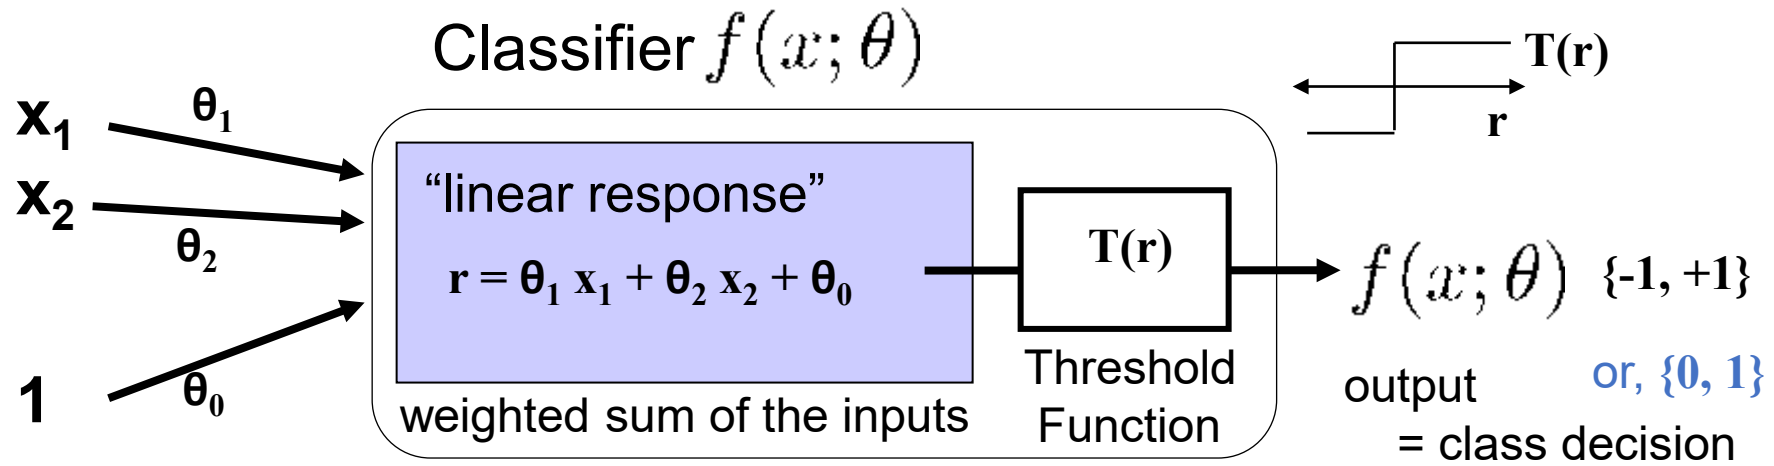
Evaluate line:

$$r = \theta_0 + \theta_1 x_1$$

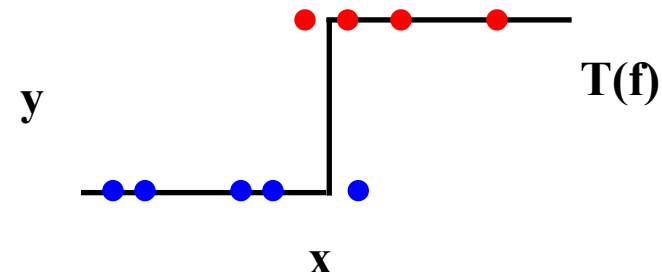
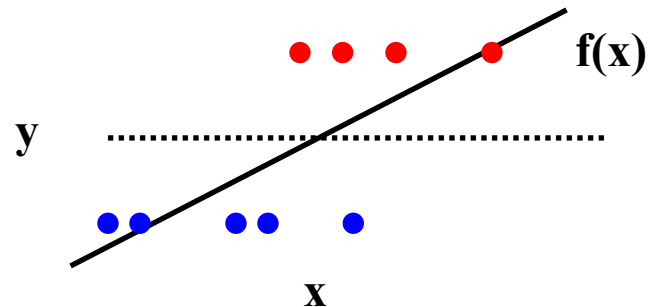
return r

- Contrast with classification
 - Classify: predict discrete-valued target y
 - Initially: “classic” binary $\{-1, +1\}$ classes; generalize later

Linear Classifier (2 features)



Visualizing for one feature “x”:

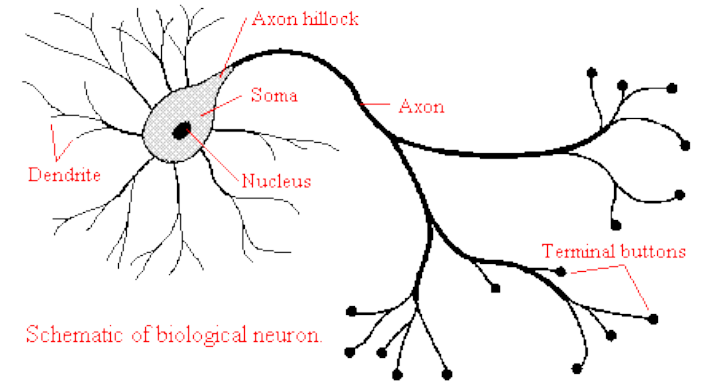


Notations

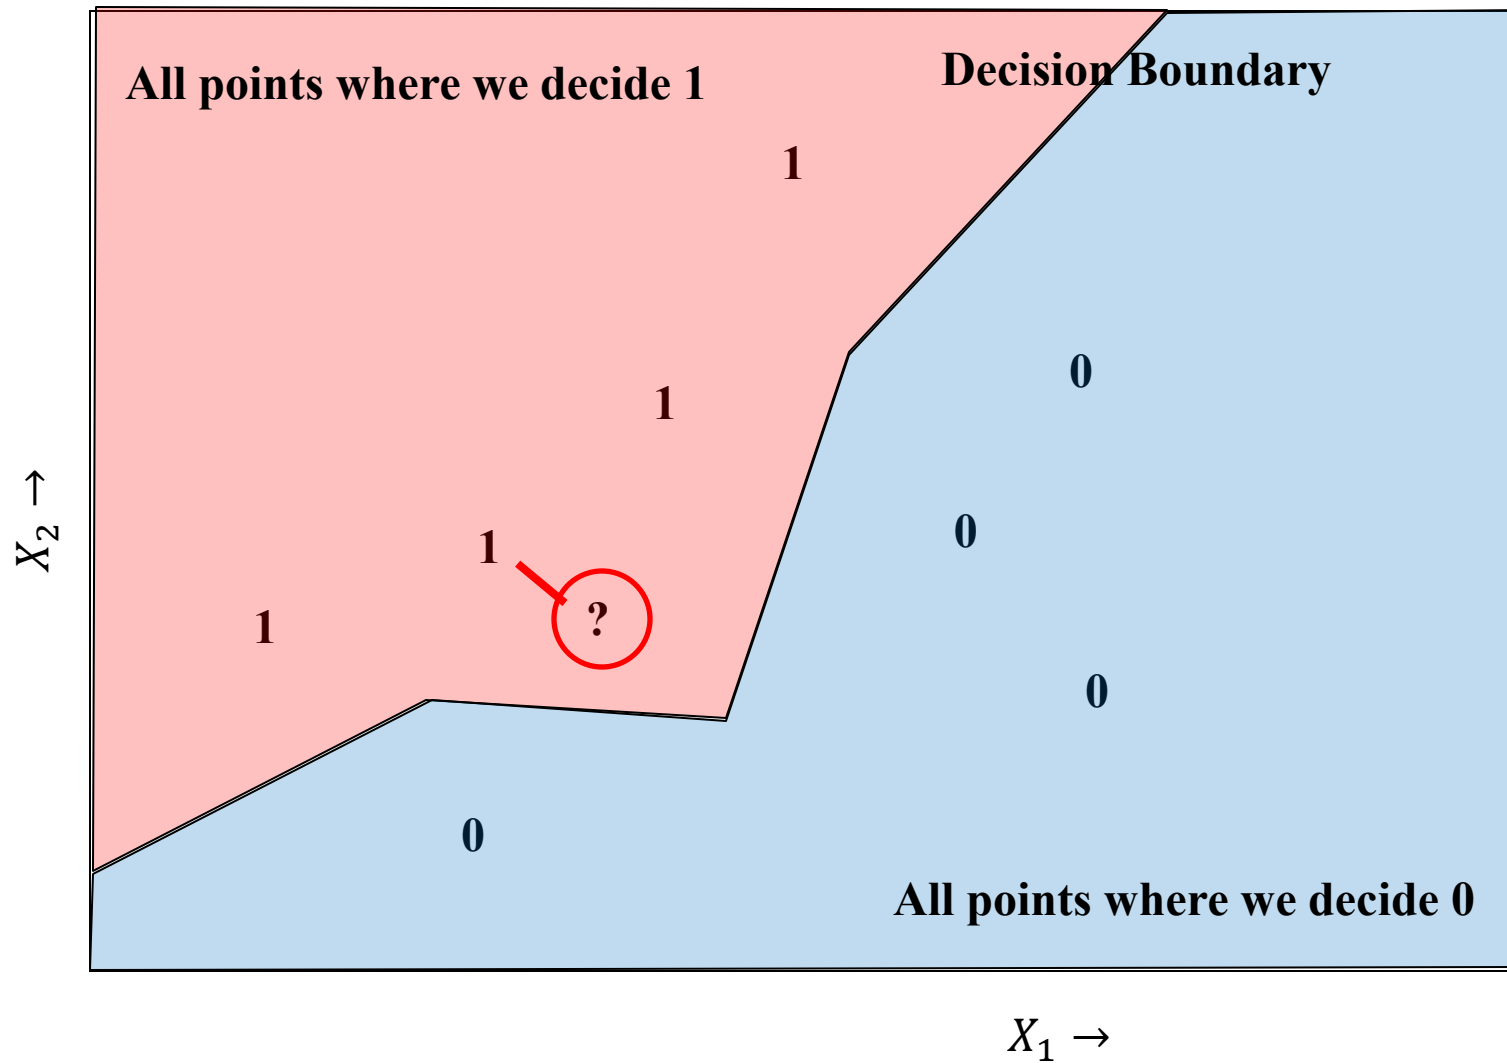
- Inputs:
 - $x_1, x_2, \dots, x_{n-1}, x_n$ are the values of the n features
 - $x_0 = 1$ (a constant input)
 - $x = (x_0, x_1, x_2, \dots, x_n)$: feature vector
- Weights (parameters):
 - $\theta_0, \theta_1, \theta_2, \dots, \theta_n$,
 - we have $n+1$ weights: one for each feature + one for the constant
 - $\theta = (\theta_0, \theta_1, \theta_2, \dots, \theta_n)$: parameter vector
- Linear response
 - $\theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n = \theta \cdot x$
- Threshold function
 - $T(r)$
- Linear classifier
 - $f(x; \theta) = T(\theta x)$

Perceptrons

- Perceptron = a linear classifier
 - The parameters θ are sometimes called weights (“w”)
 - real-valued constants (can be positive or negative)
 - Input features $x_1 \dots x_n$;
- A perceptron calculates 2 quantities:
 - 1. A weighted sum of the input features
 - 2. This sum is then thresholded by the $T(.)$ function
- Perceptron: a simple artificial model of human neurons
 - weights = “synapses”
 - threshold = “neuron firing”



Nearest neighbor classifier



Example: Gaussian Bayes for Iris Data

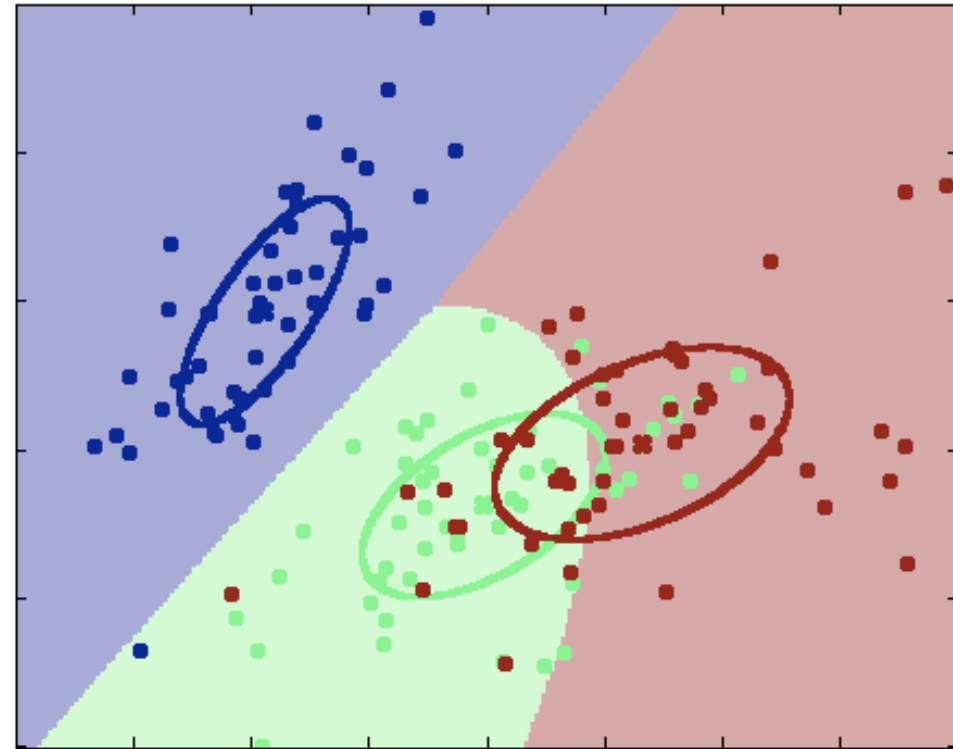
- Fit Gaussian distribution to each class $\{0,1,2\}$

$$p(y) = \text{Discrete}(\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$$

$$p(x_1, x_2 | y = 0) = \mathcal{N}(x; \mu_0, \Sigma_0)$$

$$p(x_1, x_2 | y = 1) = \mathcal{N}(x; \mu_1, \Sigma_1)$$

$$p(x_1, x_2 | y = 2) = \mathcal{N}(x; \mu_2, \Sigma_2)$$



Perceptron Decision Boundary

- The perceptron is defined by the decision algorithm:

$$f(x; \theta) = \begin{cases} +1 & \text{if } \theta \cdot x^T > 0 \\ -1 & \text{otherwise} \end{cases} \quad \text{or } f(x; \theta) = T(\theta x)$$

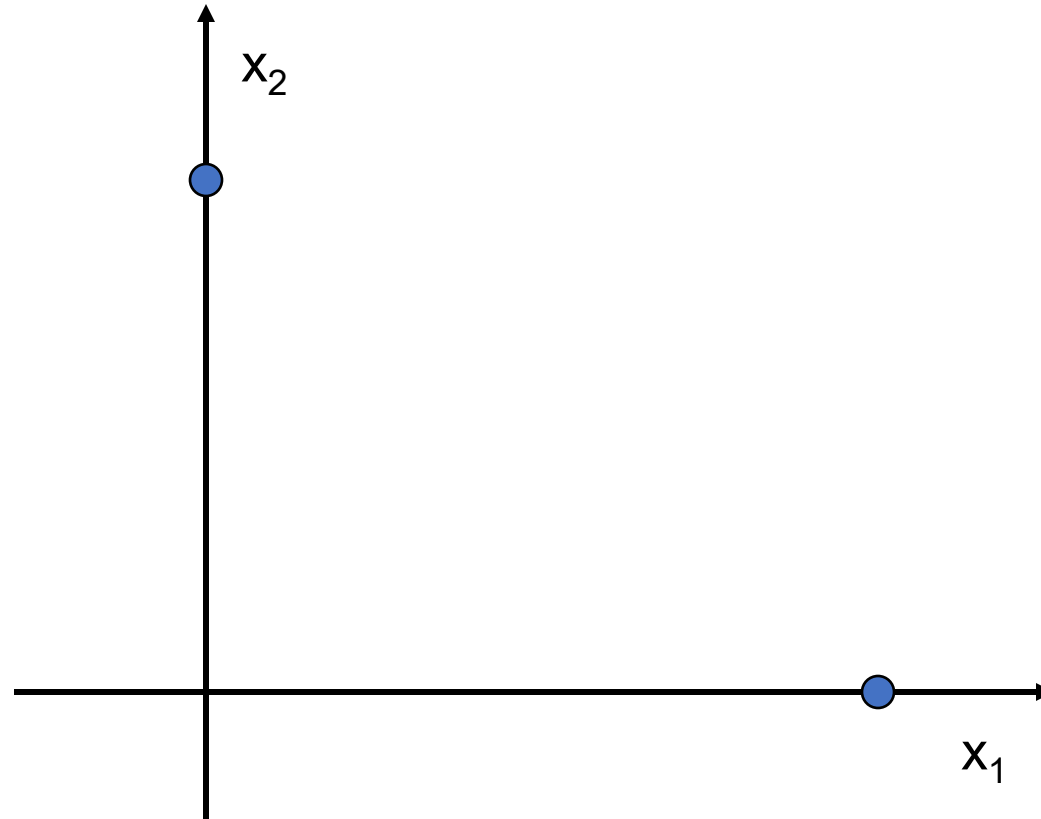
- The perceptron represents a hyperplane decision surface in n-dimensional space
 - A point in 1D, a line in 2D, a plane in 3D, etc.
- The equation of the hyperplane is given by

$$\theta \cdot \underline{x}^T = 0$$

This defines the set of points that are on the boundary.

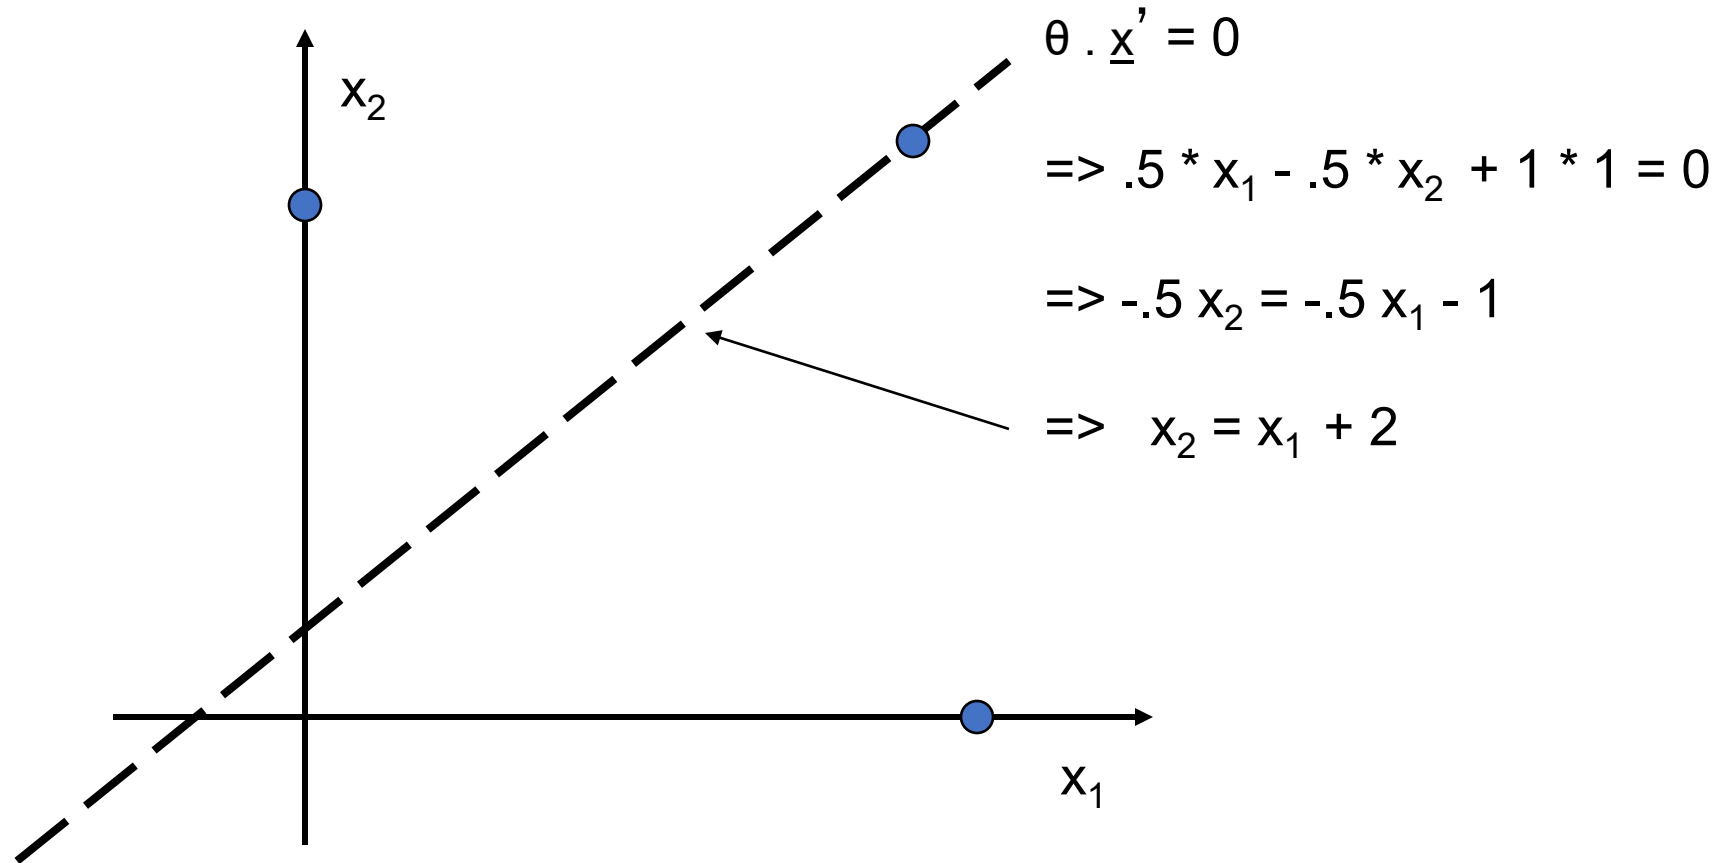
Example, Linear Decision Boundary

$$\begin{aligned}\theta &= (\theta_0, \theta_1, \theta_2) \\ &= (1, .5, -.5)\end{aligned}$$

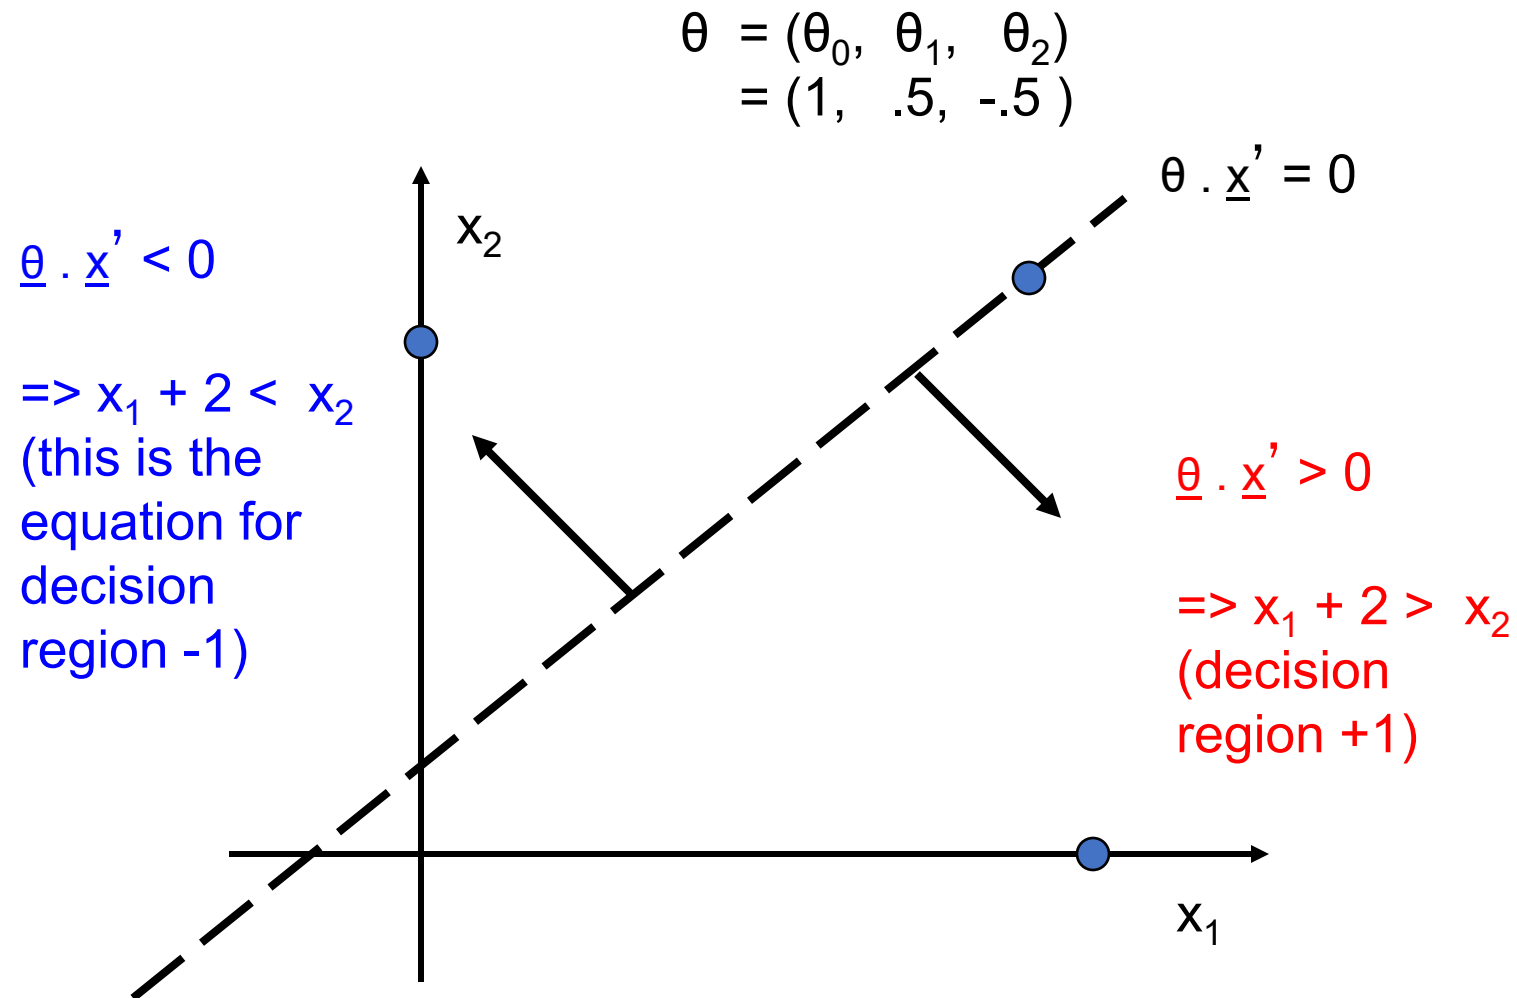


Example, Linear Decision Boundary

$$\begin{aligned}\theta &= (\theta_0, \theta_1, \theta_2) \\ &= (1, .5, -.5)\end{aligned}$$



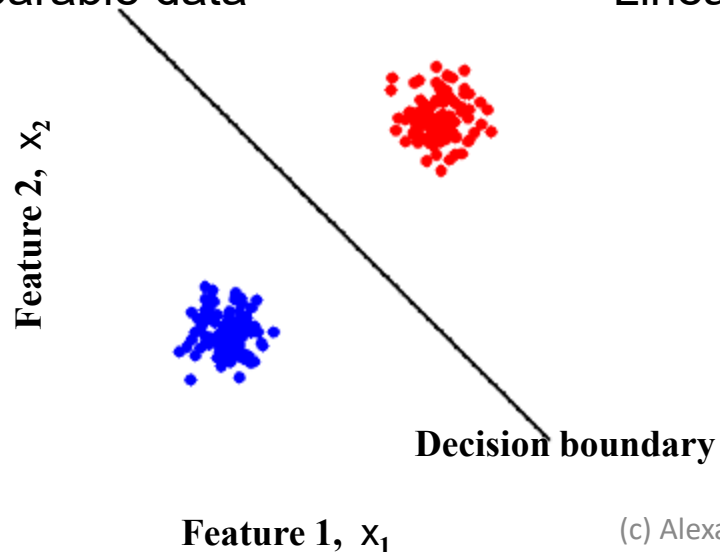
Example, Linear Decision Boundary



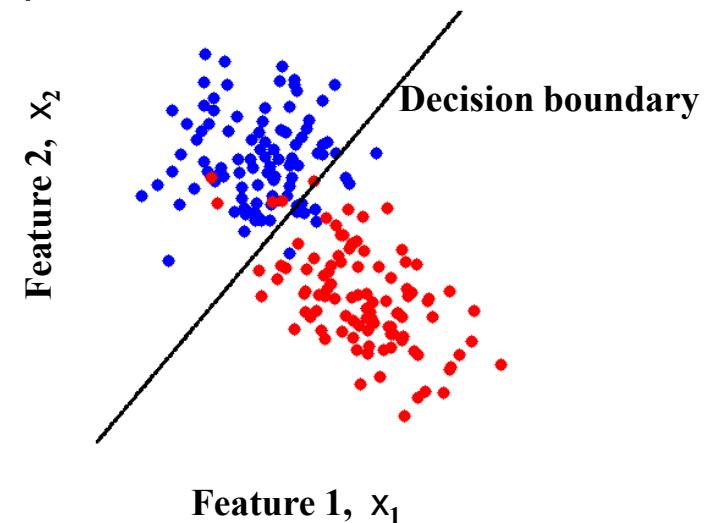
Separability

- A data set is separable by a learner if
 - There is some instance of that learner that correctly predicts all the data points
- Linearly separable data
 - Can separate the two classes using a straight line in feature space
 - in 2 dimensions the decision boundary is a straight line

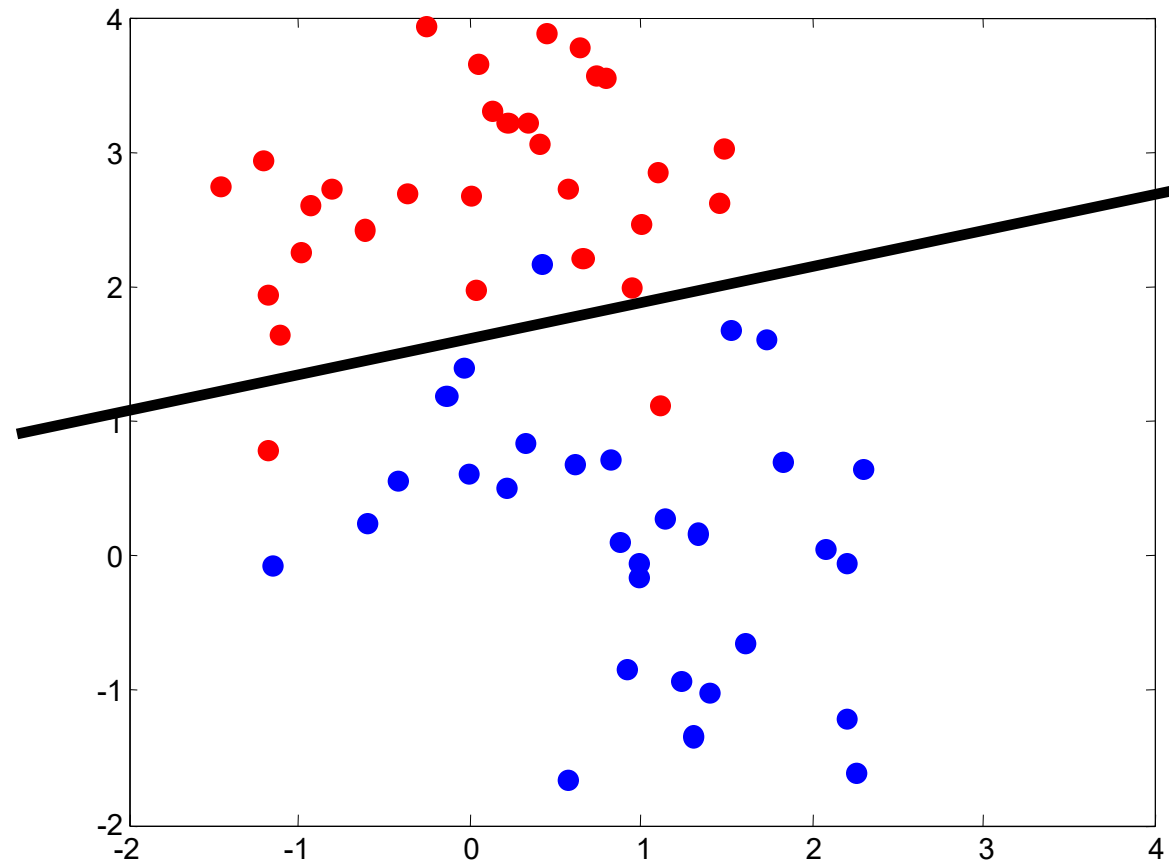
Linearly separable data



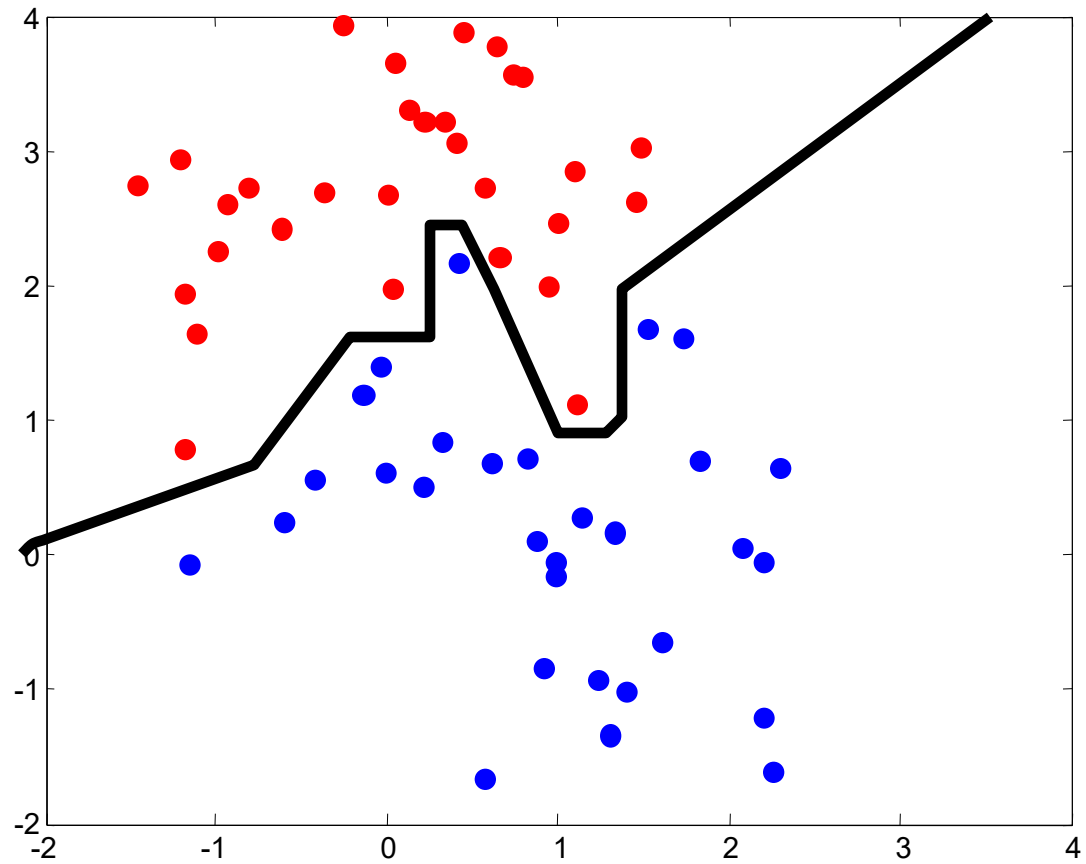
Linearly non-separable data



Another example



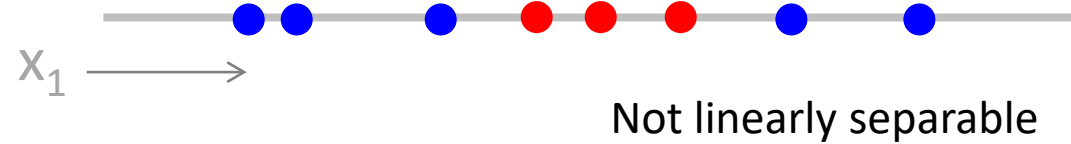
Non-linear decision boundary



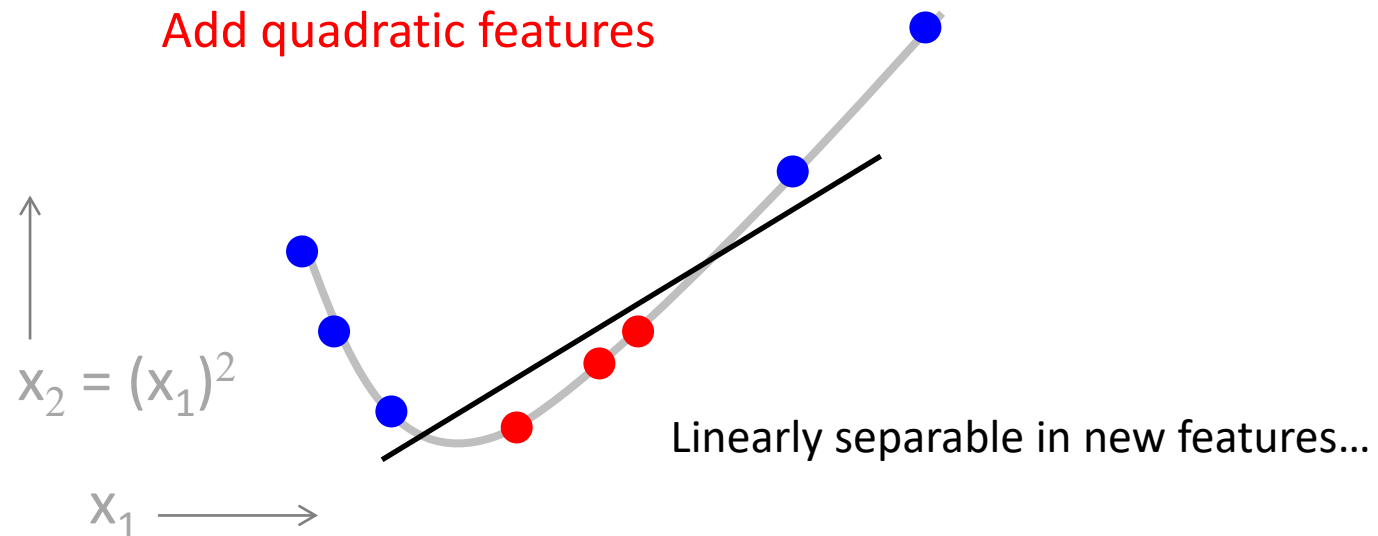
Adding features

- Linear classifier can't learn some functions

1D example:



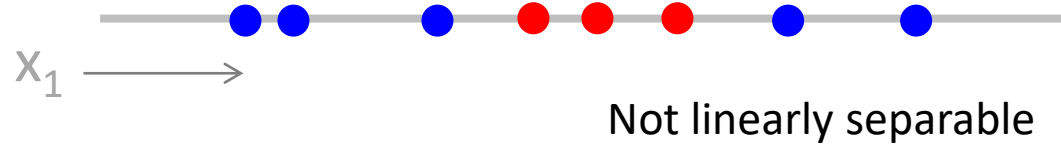
Add quadratic features



Adding features

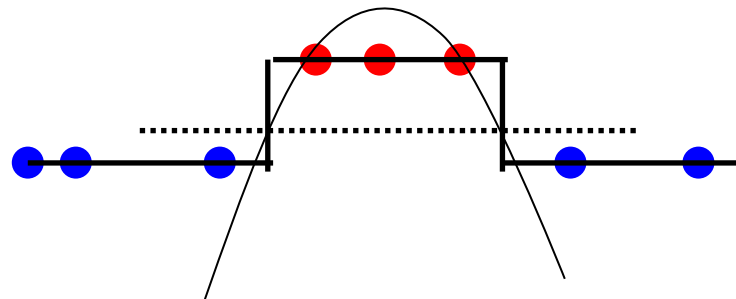
- Linear classifier can't learn some functions

1D example:



Quadratic features, visualized in original feature space:

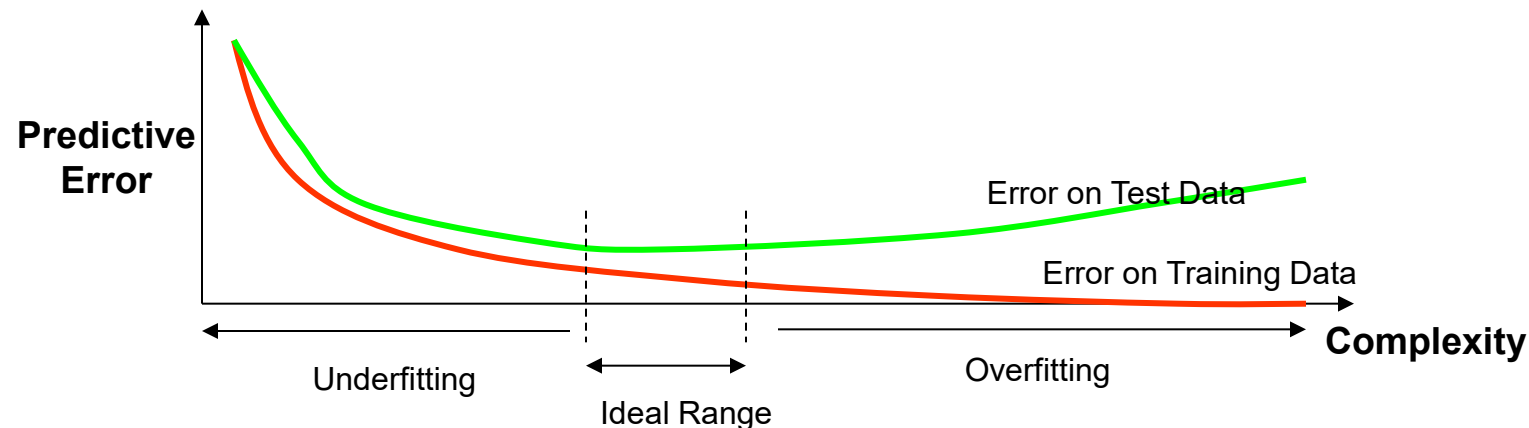
$$y = T(a x^2 + b x + c)$$



More complex decision boundary: $ax^2+bx+c = 0$

Effect of dimensionality

- Data are increasingly separable in high dimension – is this a good thing?
- “Good”
 - Separation is easier in higher dimensions (for fixed # of data m)
 - Increase the number of features, and even a linear classifier will eventually be able to separate all the training examples!
- “Bad”
 - Remember training vs. test error? Remember overfitting?
 - Increasingly complex decision boundaries can eventually get all the training data right, but it doesn't necessarily bode well for test data...



Linear Classifiers: Learning

Learning the Classifier Parameters

- Learning from Training Data:
 - training data = labeled feature vectors
 - Find parameter values that predict well (low error)
 - error is estimated on the training data
 - “true” error will be on future test data
- Define a loss function $J(\theta)$:
 - Classifier error rate (for a given set of weights $\underline{\theta}$ and labeled data)
- Minimize this loss function (or, maximize accuracy)
 - An optimization or search problem over the vector $(\theta_1, \theta_2, \theta_0)$

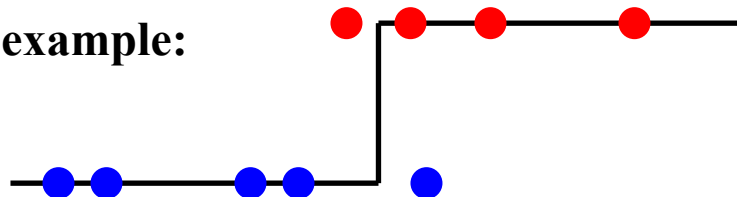
Training a linear classifier

- How should we measure error?
 - Natural measure = “fraction we get wrong” (error rate)

$$\text{err}(\theta) = \frac{1}{m} \sum_i \mathbb{1}[y^{(i)} \neq f(x^{(i)}; \theta)] \quad \text{where} \quad \mathbb{1}[y \neq \hat{y}] = \begin{cases} 1 & y \neq \hat{y} \\ 0 & \text{o.w.} \end{cases}$$

- But, hard to train via gradient descent
 - Not continuous
 - As decision boundary moves, errors change abruptly

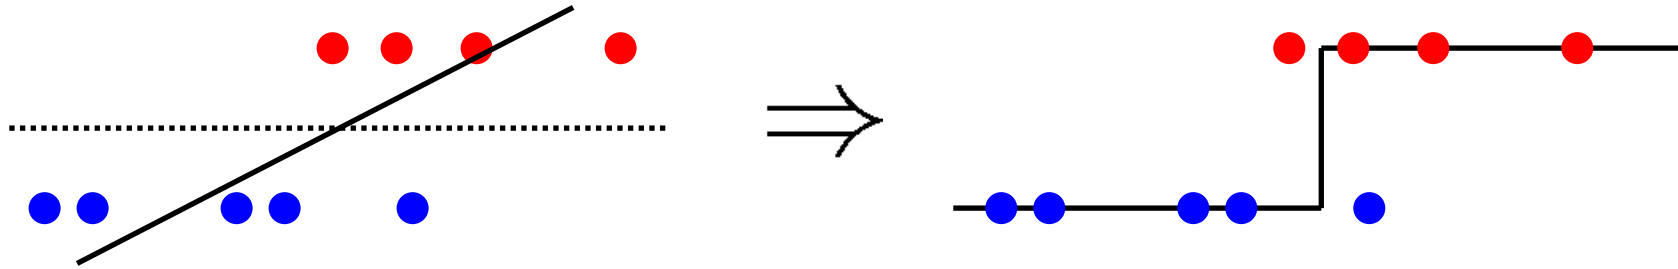
1D example:



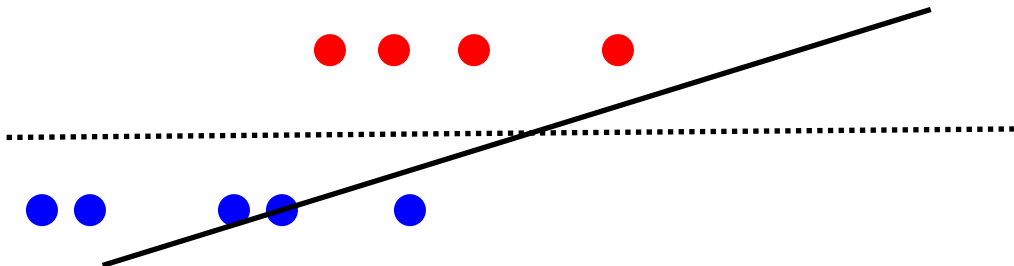
$$\begin{aligned} T(f) &= -1 & \text{if } f < 0 \\ T(f) &= +1 & \text{if } f > 0 \end{aligned}$$

Linear regression?

- Simple option: set θ using linear regression



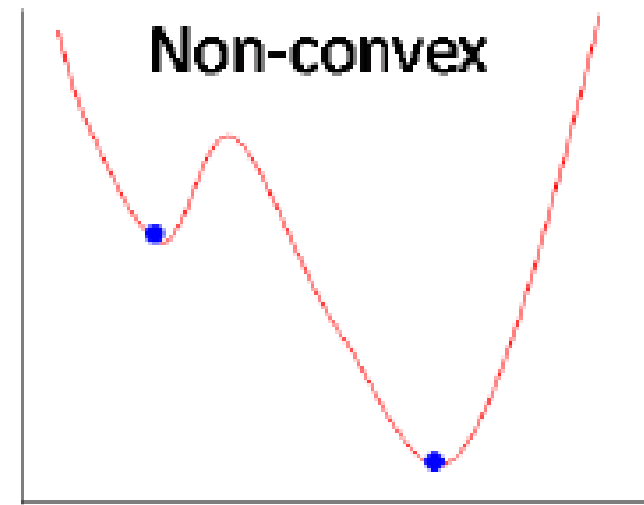
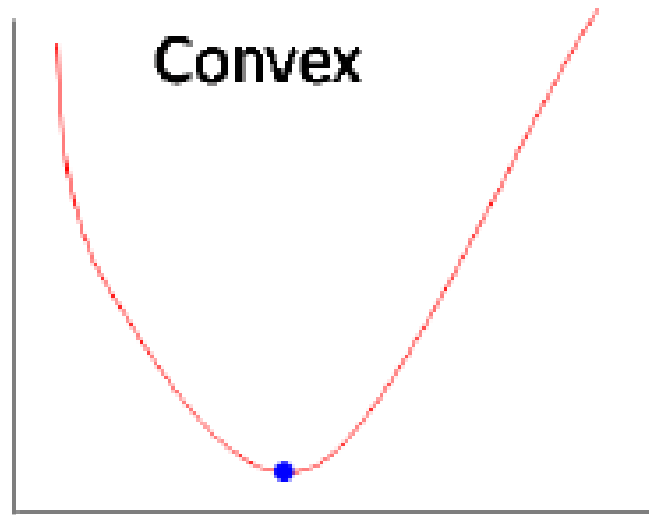
- In practice, this often doesn't work so well...
 - Consider adding a distant but “easy” point
 - MSE distorts the solution



Surrogate loss functions

- Another solution: use a “smooth” loss
 - e.g., use a smooth surrogate function to approximate the loss function

Convex?



Surrogate loss functions

Class $y = \{-1, 1\}$

$$J(\theta) = \frac{1}{m} \sum_i \mathbf{1}[y^{(i)} \neq \text{sign}(\theta x^{(i)})]$$

$$= \frac{1}{m} \sum_i \mathbf{1}[y^{(i)} \cdot \theta x^{(i)} < 0]$$

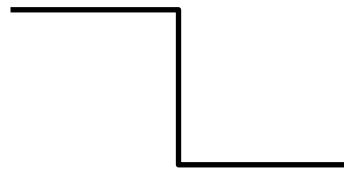
$$= \frac{1}{m} \sum_i L(y^{(i)} \cdot \theta x^{(i)})$$

0 / 1 Loss

Surrogate loss functions

Class $y = \{0, 1\}$

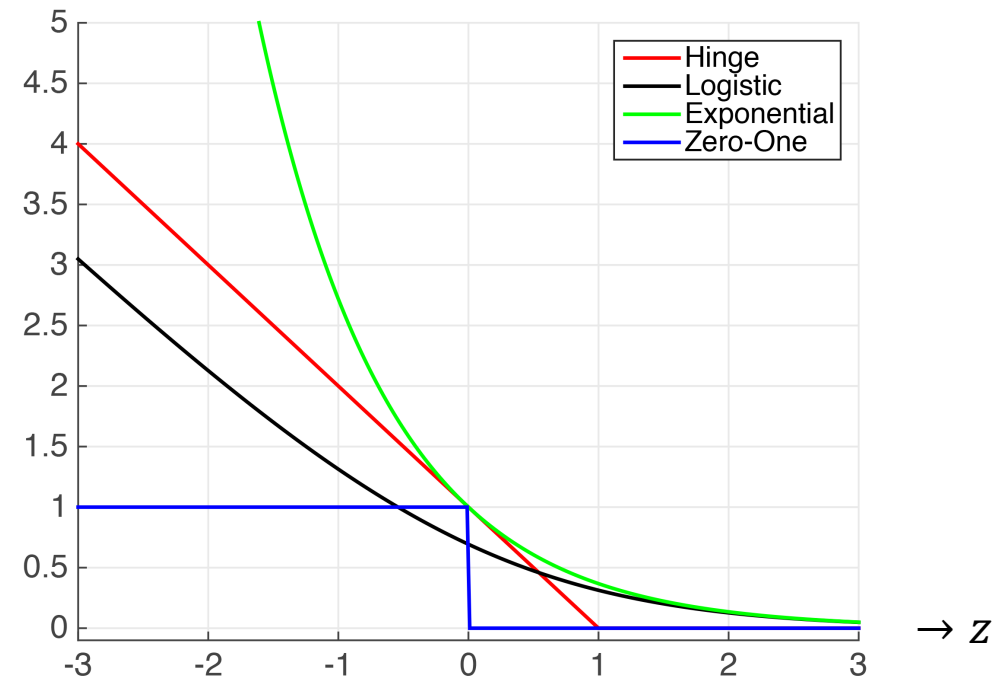
$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_i \mathbf{1} [y^{(i)} \neq \mathbf{1} [\theta x^{(i)} > 0]] \\ &= \frac{1}{m} \sum_i (y^{(i)} \mathbf{1} [\theta x^{(i)} < 0] + (1 - y^{(i)}) \mathbf{1} [\theta x^{(i)} > 0]) \\ &= \frac{1}{m} \sum_i (y^{(i)} L(\theta x^{(i)}) + (1 - y^{(i)}) L(-\theta x^{(i)})) \end{aligned}$$



0 / 1 Loss

Surrogate loss functions

- 0-1: $L(z) = \mathbf{1}[z < 0]$
- Logistic: $L(z) = -\log \sigma(z) = -\log \frac{1}{1+e^{-z}}$
- Exponential: $L(z) = e^{-\beta z}$
- Hinge: $L(z) = \max\{0, 1 - z\}$
- ...



Generic classification formulation

Class $y = \{0, 1\}$

$$J(\theta) = \frac{1}{m} \sum_i \left(y^{(i)} \phi(\theta x^{(i)}) + (1 - y^{(i)}) \phi(-\theta x^{(i)}) \right)$$

Class $y = \{-1, 1\}$

$$J(\theta) = \frac{1}{m} \sum_i \phi(y^{(i)} \theta x^{(i)})$$

Generic classification formulation w/ Regularization

Class $y = \{0, 1\}$

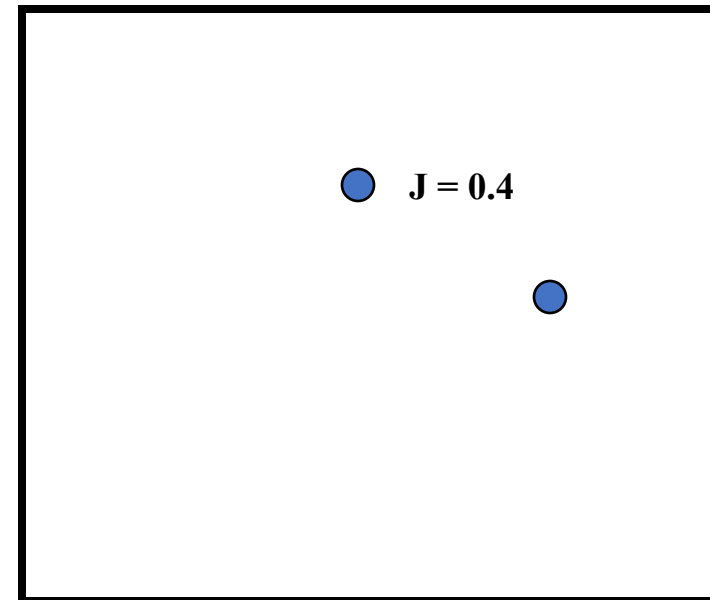
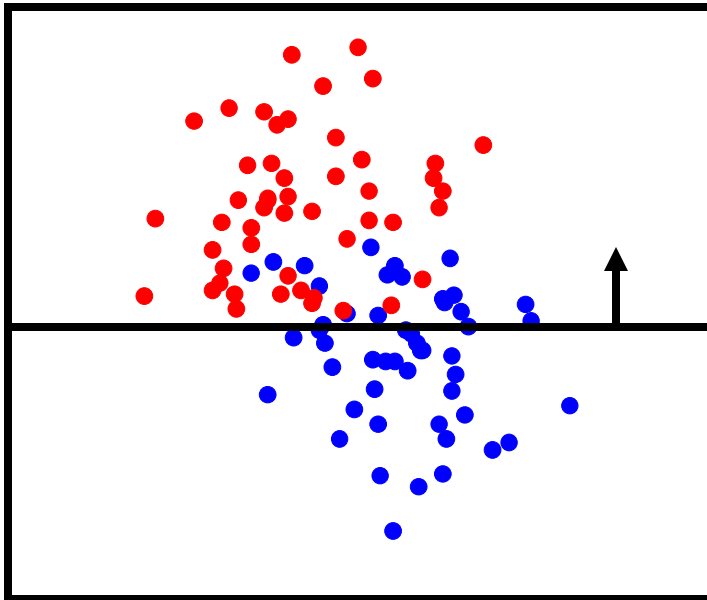
$$J(\theta) = \frac{1}{m} \sum_i \left(y^{(i)} \phi(\theta x^{(i)}) + (1 - y^{(i)}) \phi(-\theta x^{(i)}) \right) + \frac{\lambda}{2m} ||\theta||^2$$

Class $y = \{-1, 1\}$

$$J(\theta) = \frac{1}{m} \sum_i \phi(y^{(i)} \theta x^{(i)}) + \frac{\lambda}{2m} ||\theta||^2$$

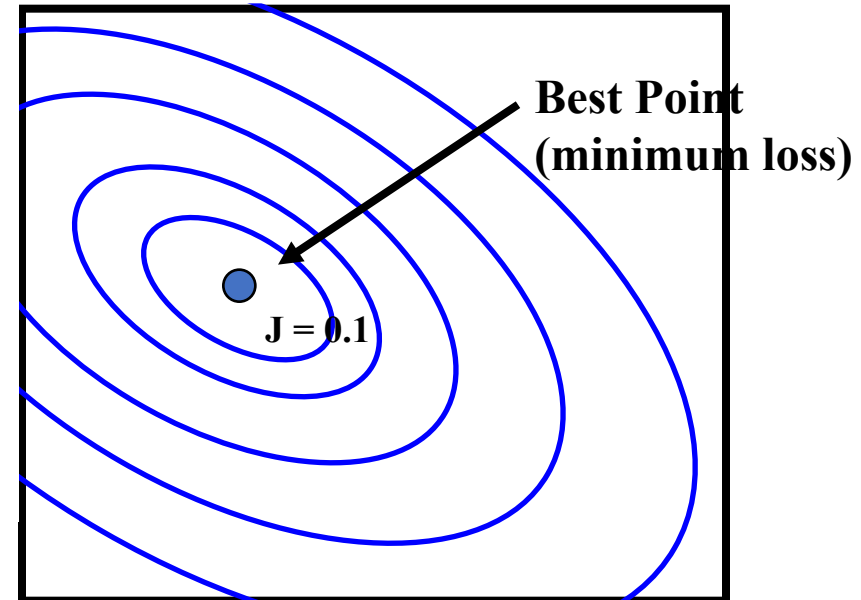
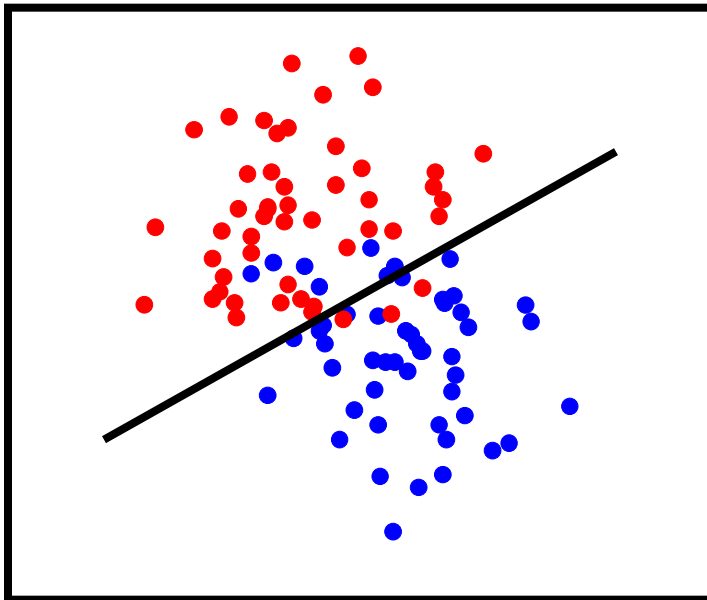
Training the Classifier

- Once we have a smooth measure of quality, we can find the “best” settings for the parameters
- Example: 2D feature space \Leftrightarrow parameter space



Training the Classifier

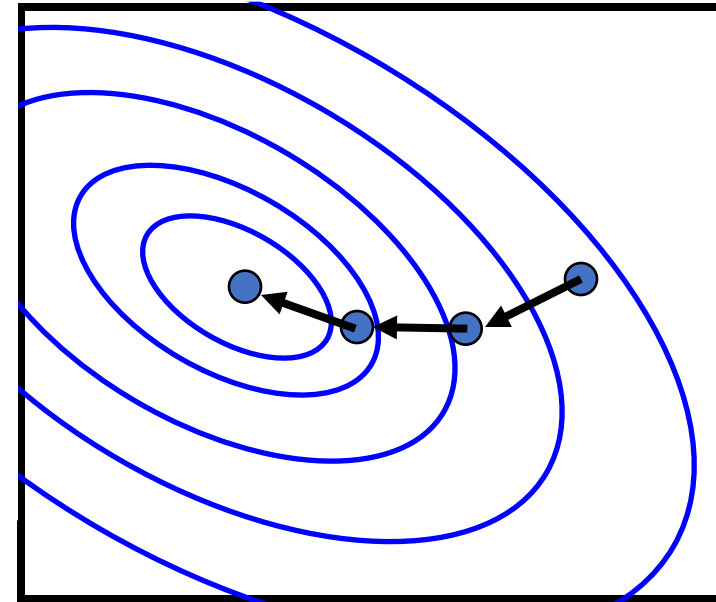
- Once we have a smooth measure of quality, we can find the “best” settings for the parameters
- Example: 2D feature space \Leftrightarrow parameter space



Minimizing the loss function

- As in linear regression, this is now just optimization
- Methods:
 - Gradient descent
 - Improve loss by small changes in parameters (“small” = learning rate)

Gradient Descent



Gradient of general loss functions

Class $y = \{0, 1\}$

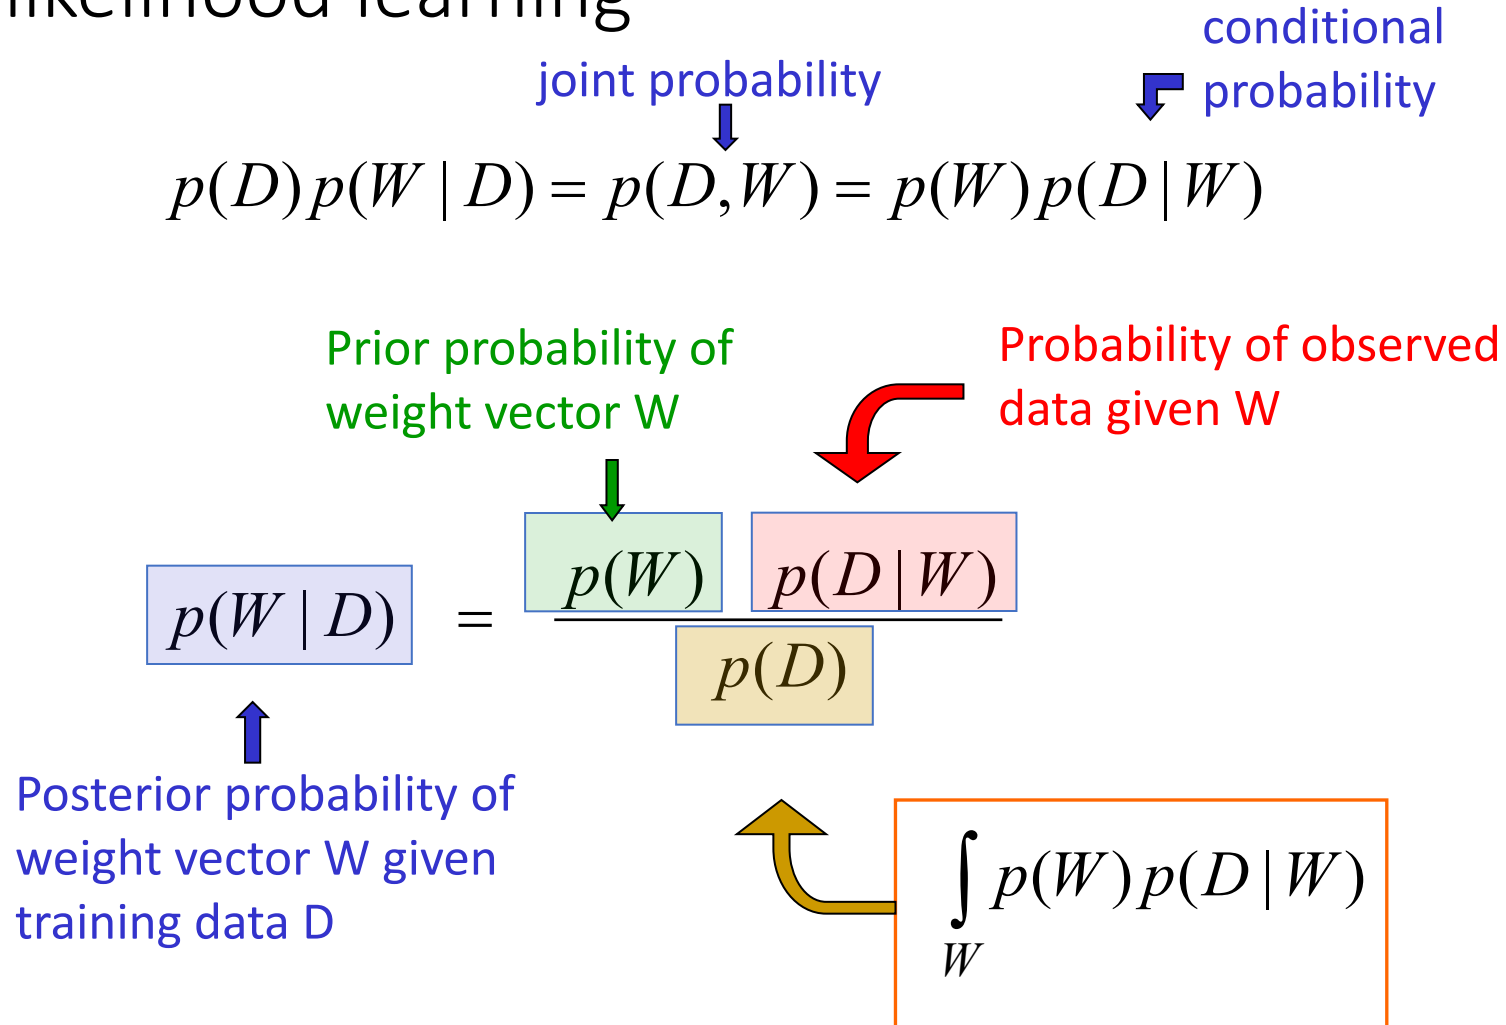
$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_i \left(y^{(i)} \frac{\partial \phi(\theta x^{(i)})}{\partial \theta_j} + (1 - y^{(i)}) \frac{\partial \phi(-\theta x^{(i)})}{\partial \theta_j} \right)$$

Class $y = \{-1, 1\}$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_i \frac{\partial \phi(y^{(i)} \theta x^{(i)})}{\partial \theta_j}$$

Logistic Regression

Maximum likelihood learning



Maximize sums of log probs

- We want to maximize the **product** of the probabilities of the outputs on the training cases
 - Assume the output errors on different training cases, i , are independent.

$$p(D|W) = \prod_i p(d^{(i)} | W)$$

- Because the log function is monotonic, it does not change where the maxima are. So we can maximize **sums** of log probabilities

$$\log p(D|W) = \sum_i \log p(d^{(i)} | W)$$

- This is called **maximum likelihood** learning.

Minimum negative log-likelihood: $-\log p(D|W) = -\sum_i \log p(d^{(i)} | W)$

For classification...

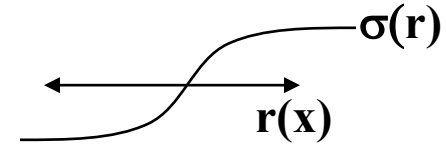
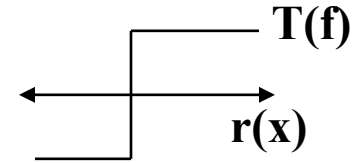
- Minimum negative log-likelihood:

$$-\log p(Y|X, \theta) = -\sum_i \log p(y^{(i)} | x^{(i)}, \theta)$$

Logistic regression

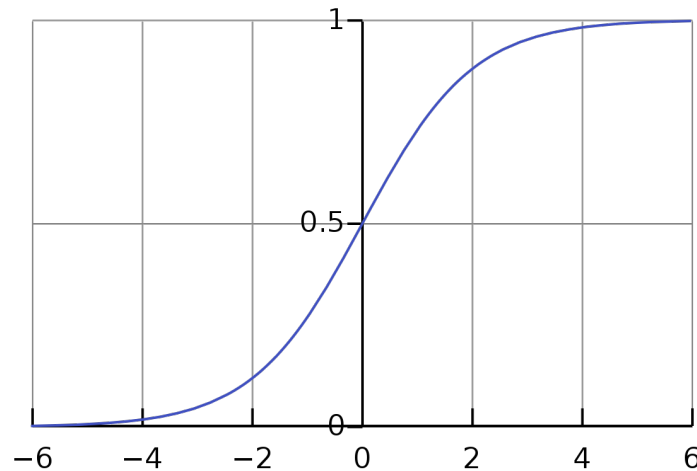
- Use a “smooth” function to approximate the threshold function

$$T(r) \Rightarrow \sigma(r)$$



- Logistic “sigmoid”, looks like an “S”

$$\sigma(r) = \frac{1}{1 + e^{-r}}$$



Logistic regression

- Interpret $\sigma(\theta x)$ as a probability that $y = 1$, i.e., $P(Y = 1|x; \theta) = \sigma(\theta x)$
- Use a negative log-likelihood loss function
 - If $y = 1$, loss is $-\log P[y = 1] = -\log \sigma(\theta x)$
 - If $y = 0$, loss is $-\log P[y = 0] = -\log(1 - \sigma(\theta x))$
- Can write this succinctly:

$$J(\underline{\theta}) = -\frac{1}{m} \left(\sum_i \underbrace{y^{(i)} \log \sigma(\theta \cdot x^{(i)})}_{\text{Nonzero only if } y=1} + \underbrace{(1-y^{(i)}) \log(1-\sigma(\theta \cdot x^{(i)}))}_{\text{Nonzero only if } y=0} \right)$$

Logistic regression

$$J(\underline{\theta}) = -\frac{1}{m} \left(\sum_i y^{(i)} \log \sigma(\theta \cdot x^{(i)}) + (1 - y^{(i)}) \log(1 - \sigma(\theta \cdot x^{(i)})) \right)$$

$$J(\theta) = \frac{1}{m} \sum_i \left(y^{(i)} \phi(\theta x^{(i)}) + (1 - y^{(i)}) \phi(-\theta x^{(i)}) \right)$$

$$\phi(z) = -\log(\sigma(z)) = -\log \frac{1}{1 + e^{-z}}$$

$$\phi(-z) = -\log(\sigma(-z)) = -\log \frac{1}{1 + e^z} = -\log \frac{e^{-z}}{1 + e^{-z}} = -\log \left(1 - \frac{1}{1 + e^{-z}} \right) = -\log(1 - \sigma(z))$$

Gradient Equations

$$(\ln z)' = \frac{1}{z}$$

$$(\sigma(z))' = \sigma(z)(1 - \sigma(z))$$

- Logistic neg-log likelihood loss:

$$J(\underline{\theta}) = -\frac{1}{m} \left(\sum_i y^{(i)} \log \sigma(\theta \cdot x^{(i)}) + (1 - y^{(i)}) \log(1 - \sigma(\theta \cdot x^{(i)})) \right)$$

- What's the derivative with respect to one of the parameters?

$$\frac{\partial J(\theta)}{\partial \theta_j}$$

$$\begin{aligned} &= -\frac{1}{m} \sum_i \left(y^{(i)} (1 - \sigma(\theta x^{(i)})) x_j^{(i)} - (1 - y^{(i)}) \sigma(\theta x^{(i)}) x_j^{(i)} \right) \\ &= \frac{1}{m} \sum_i (\sigma(\theta x^{(i)}) - y^{(i)}) x_j^{(i)} \end{aligned}$$

(Batch) Gradient descent

```
Initialize  $\theta$   
Do {  
     $\theta \leftarrow \theta - \alpha \nabla J(\theta)$   
} while (stop condition)
```

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \left(y^{(i)} \log \sigma(\theta x^{(i)}) + (1 - y^{(i)}) \log (1 - \sigma(\theta x^{(i)})) \right)$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m (\sigma(\theta x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Stochastic gradient descent

- Instead of evaluating gradient over all examples evaluate it for each **individual** training example

```
Initialize  $\theta$   
Do {  
  for each  $i$   
     $\theta \leftarrow \theta - \alpha \nabla J^{(i)}(\theta)$   
} while (stop condition)
```

$$J^{(i)}(\theta) = y^{(i)} \log \sigma(\theta x^{(i)}) + (1 - y^{(i)}) \log (1 - \sigma(\theta x^{(i)}))$$

$$\frac{\partial J^{(i)}(\theta)}{\partial \theta_j} = (\sigma(\theta x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Stochastic gradient descent

- Update based on each datum at a time
 - Find residual and the gradient of its part of the error & update

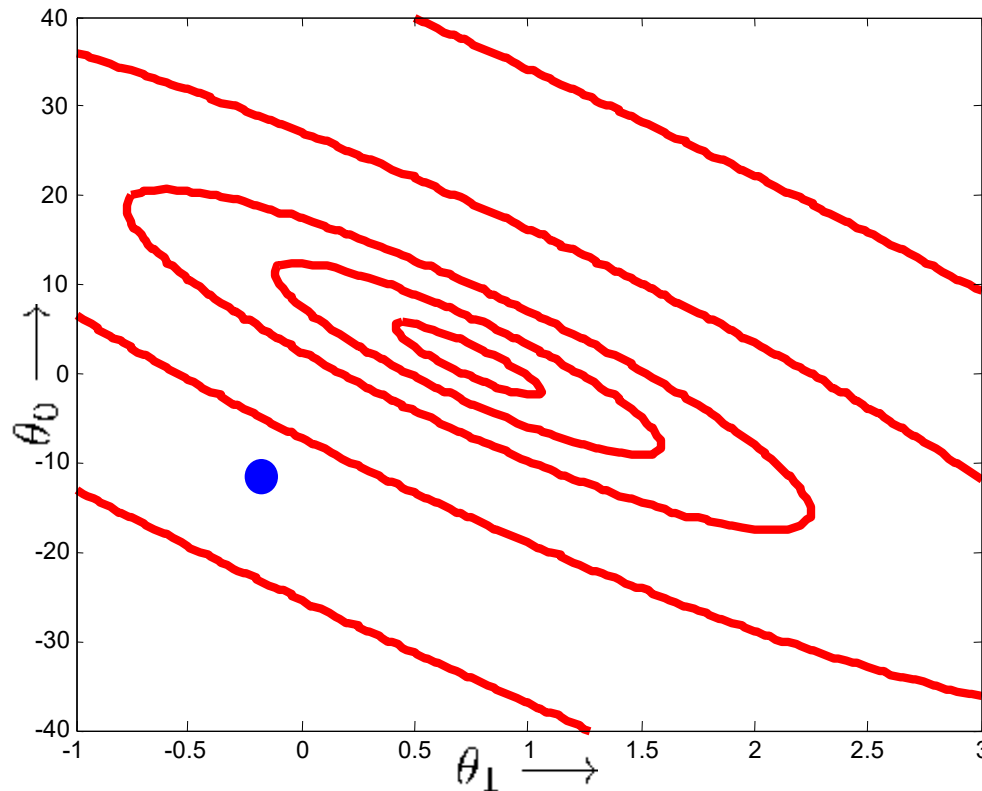
Initialize θ

Do {

 for $i = 1 : m$

$\theta \leftarrow \theta - \alpha \nabla J^{(i)}(\theta)$

 } while (stop condition)



Stochastic gradient descent

- Update based on each datum at a time
 - Find residual and the gradient of its part of the error & update

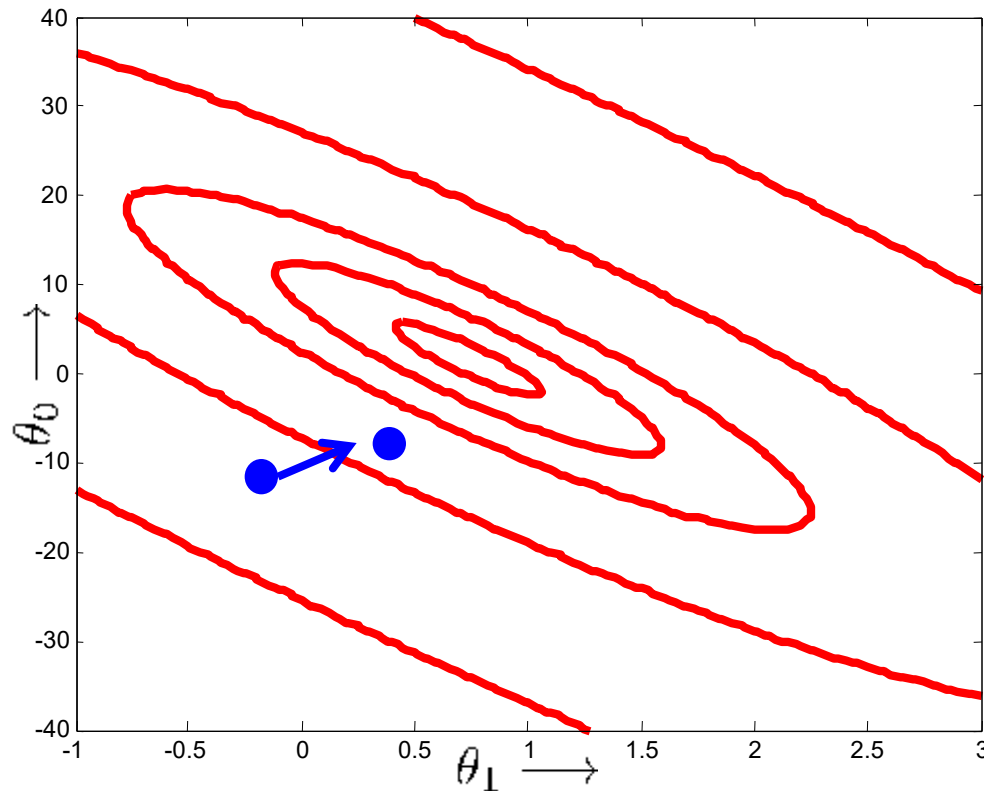
Initialize θ

Do {

 for $i = 1 : m$

$\theta \leftarrow \theta - \alpha \nabla J^{(i)}(\theta)$

 } while (stop condition)



Stochastic gradient descent

- Update based on each datum at a time
 - Find residual and the gradient of its part of the error & update

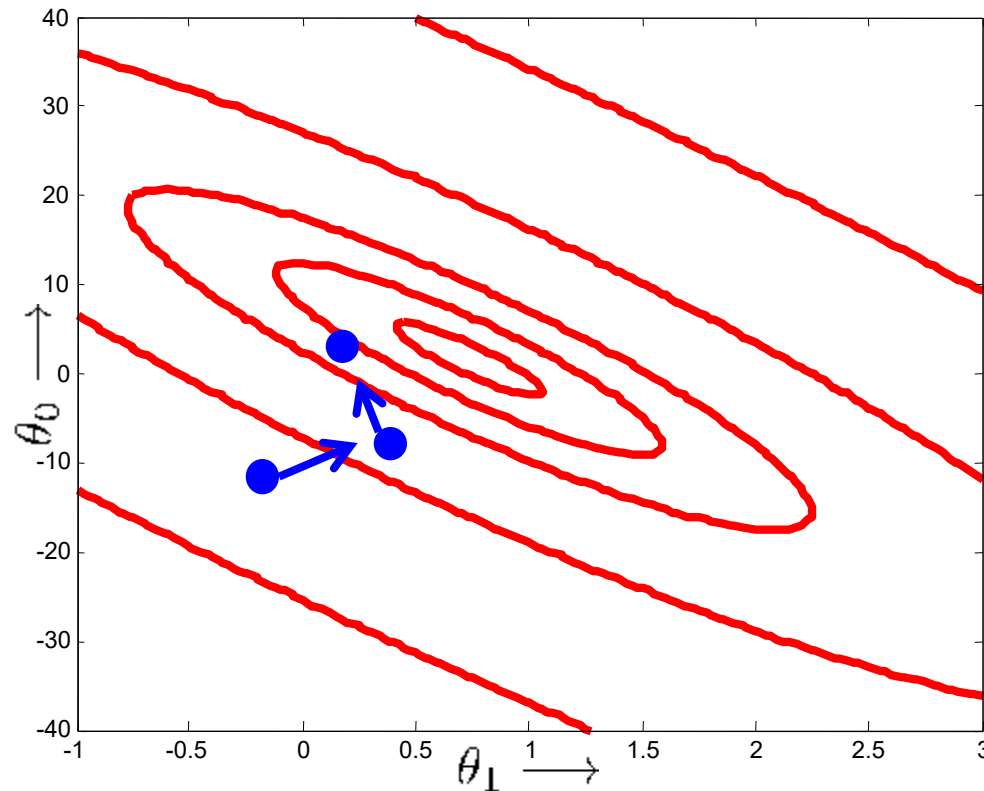
Initialize θ

Do {

 for $i = 1 : m$

$\theta \leftarrow \theta - \alpha \nabla J^{(i)}(\theta)$

 } while (stop condition)



Stochastic gradient descent

- Update based on each datum at a time
 - Find residual and the gradient of its part of the error & update

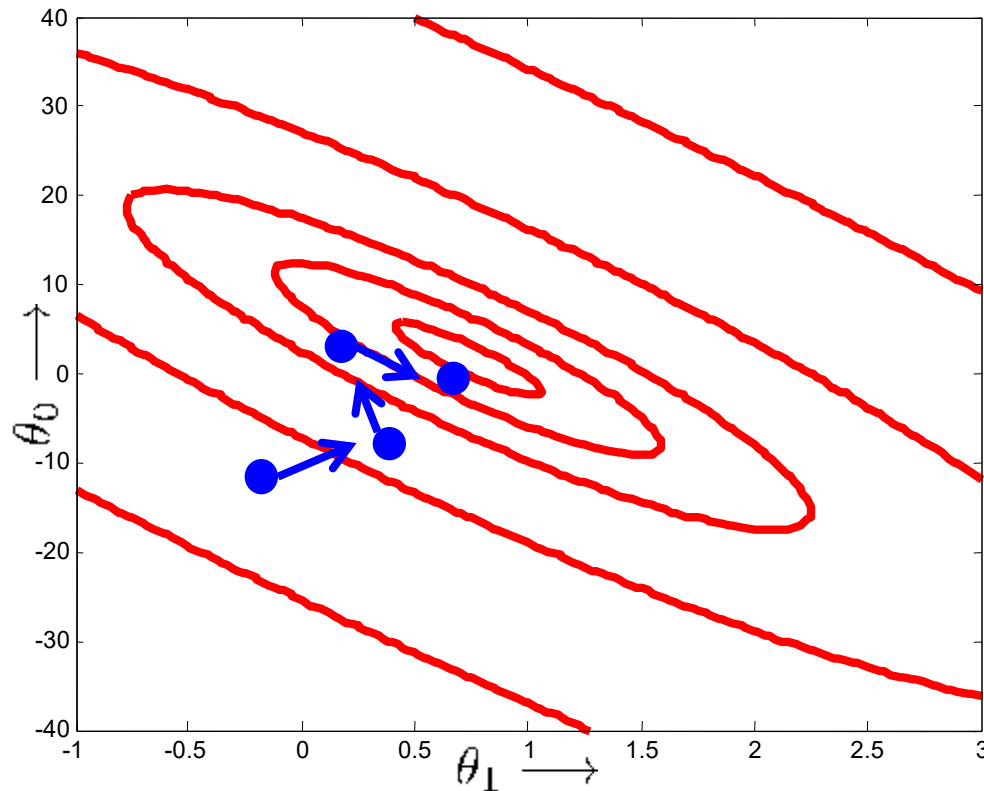
Initialize θ

Do {

 for $i = 1 : m$

$\theta \leftarrow \theta - \alpha \nabla J^{(i)}(\theta)$

 } while (stop condition)



Stochastic gradient descent

- Benefits
 - Lots of data = many more updates per pass
 - Computationally faster
- Drawbacks
 - No longer strictly “descent”
 - Stopping conditions may be harder to evaluate
(Can use “running estimates” of $J(\cdot)$, etc.)
- Related: mini-batch updates, etc.

```
Initialize  $\theta$ 
Do {
  for  $i = 1 : m$ 
     $\theta \leftarrow \theta - \alpha \nabla J^{(i)}(\theta)$ 
} while (stop condition)
```

Summary

- Linear classifier \Leftrightarrow perceptron
- Measuring quality of a decision boundary
 - Error rate (0/1 loss)
 - Surrogate functions
- Learning the weights of a linear classifier from data
 - Reduces to an optimization problem
 - Perceptron algorithm
 - Using surrogate functions, we can do gradient descent
 - Gradient equations & update rules (BGD and SGD)
 - Multiclass logistic regression (softmax function)