


t-Distributed Stochastic Neighbour Embedding (t-SNE)

Adopted from slides by Ethan Fetaya, James Lucas and Emad Andrews at University of Toronto


Local embedding

- t-SNE is an alternative dimensionality reduction algorithm.
- PCA tries to find a **global** structure
 - Low dimensional subspace
 - Can lead to local inconsistencies
 - Far away point can become nearest neighbors
- t-SNE tries to preserve **local** structure
 - Low dimensional neighborhood should be the same as original neighborhood.
 - Unlike PCA, t-SNE almost only used for visualization
 - No easy way to embed new points

PCA 2 dimensions embedding for MNIST



t-SNE 2 dimensions embedding for MNIST



Stochastic Neighbor Embedding (SNE)

- SNE basic idea:
 - “Encode” high dimensional neighborhood information as a distribution Intuition: Random walk between data points.
 - High probability to jump to a close point
 - Find low dimensional points such that their neighborhood distribution is similar.
 - How do you measure distance between distributions?
 - Most common measure: KL divergence

Neighborhood Distributions

- Consider the neighborhood around an input data point $\mathbf{x}_i \in \mathbb{R}^d$
- Imagine that we have a Gaussian distribution centered around \mathbf{x}_i
- Then the probability that \mathbf{x}_i chooses some other datapoint \mathbf{x}_j as its neighbor is in proportion with the density under this Gaussian
- A point closer to \mathbf{x}_i will be more likely than one further away

Probabilities P_{ij}

- The $i \rightarrow j$ probability is the probability that point \mathbf{x}_i chooses \mathbf{x}_j as its neighbor

$$P_{j|i} = \frac{\exp(-\|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-\|\mathbf{x}^{(i)} - \mathbf{x}^{(k)}\|^2 / 2\sigma_i^2)}$$


- With $P_{i|i} = 0$
 - The parameter σ_i sets the size of the neighborhood
 - Very low σ_i - all the probability is in the nearest neighbor.
 - Very high σ_i - Uniform weights.
 - Here we set σ_i **differently** for each data point
 - Results depend heavily on σ_i - it defines the neighborhoods we are trying to preserve.
 - Final distribution over pairs is symmetrized: $P_{ij} = 1/2N(P_{i|j} + P_{j|i})$

Perplexity

- For each distribution P_i (depends on σ_i) we define the perplexity
 - $perp(P_i) = 2^{H(P_i)}$ where $H(P) = -\sum_j P_{j|i} \log(P_{j|i})$ is the entropy.
- If P is uniform over k elements - perplexity is k .
 - Smooth version of k in kNN
 - Low perplexity = small σ
 - High perplexity = large σ
- Define the desired perplexity and set σ_i to get that (binary search)
- Values between 5-50 usually work well
- Important parameter - different perplexity can capture different scales in the data

t-SNE Practical Examples

Perplexity Settings Matter



Original


Perplexity = 2

Perplexity = 30

Perplexity = 100

t-SNE Practical Examples

Cluster Sizes are Meaningless



Original

Perplexity = 5

Perplexity = 50

SNE objective

- Given $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)} \in \mathbb{R}^D$ we define the distribution P_{ij}
- Goal: Find good embedding $\mathbf{y}^{(1)}, \dots, \mathbf{y}^{(N)} \in \mathbb{R}^d$ for some $d < D$ (normally 2 or 3)
- How do we measure an embedding quality?
- For points $\mathbf{y}^{(1)}, \dots, \mathbf{y}^{(N)} \in \mathbb{R}^d$ we can define distribution Q similarly the same

$$Q_{ij} = \frac{\exp(-\|\mathbf{y}^{(i)} - \mathbf{y}^{(j)}\|^2)}{\sum_k \sum_{l \neq k} \exp(-\|\mathbf{y}^{(l)} - \mathbf{y}^{(k)}\|^2)}$$

- Optimize Q to be close to P
 - Minimize KL-divergence
- The embeddings $\mathbf{y}^{(1)}, \dots, \mathbf{y}^{(N)} \in \mathbb{R}^d$ are the parameters we are optimizing

SNE algorithm

- We have P , and are looking for $\mathbf{y}^{(1)}, \dots, \mathbf{y}^{(N)} \in R^d$ such that the distribution Q we infer will minimize $L(Q) = KL(P||Q)$.
- Note that $KL(P||Q) = \sum_{ij} P_{ij} \log\left(\frac{P_{ij}}{Q_{ij}}\right) \propto -\sum_{ij} P_{ij} \log(Q_{ij})$
- Can show that $\frac{\partial L}{\partial \mathbf{y}^{(i)}} = \sum_j (P_{ij} - Q_{ij})(\mathbf{y}^{(i)} - \mathbf{y}^{(j)})$
- Main issue - crowding problem.

Crowding Problem

- In high dimension we have more room, points can have a lot of different neighbors
- In 2D a point can have a few neighbors at distance one all far from each other - what happens when we embed in 1D?
- This is the “crowding problem” - we don’t have enough room to accommodate all neighbors.
- This is one of the biggest problems with SNE.
- t-SNE solution: Change the Gaussian in Q to a heavy tailed distribution.
 - if Q changes slower, we have more “wiggle room” to place points at.

t-SNE


- t-Distributed Stochastic Neighbor Embedding
 - Student-t Probability density $p(x) \propto (1 + \frac{x^2}{\nu})^{-(\nu+1)/2}$
- Probability goes to zero much slower than a Gaussian.
- Can show it is equivalent to averaging Gaussians with some prior over σ
- We can now redefine Q_{ij} as

$$Q_{ij} = \frac{(1 + \|\mathbf{y}_i - \mathbf{y}_j\|^2)^{-1}}{\sum_k \sum_{l \neq k} (1 + \|\mathbf{y}_k - \mathbf{y}_l\|^2)^{-1}}$$

- We leave P_{ij} as is

t-SNE

Blue = Gaussian
Red = Student's t



t-SNE gradients

- Can show that the gradients of t-SNE objective are

$$\frac{\partial L}{\partial \mathbf{y}^{(i)}} = \sum_j (P_{ij} - Q_{ij})(\mathbf{y}^{(i)} - \mathbf{y}^{(j)})(1 + \|\mathbf{y}_i - \mathbf{y}_j\|^2)^{-1}$$

- Compare to the SNE gradients

$$\frac{\partial L}{\partial \mathbf{y}^{(i)}} = \sum_j (P_{ij} - Q_{ij})(\mathbf{y}^{(i)} - \mathbf{y}^{(j)})$$

Algorithm

Algorithm 1: Simple version of t-Distributed Stochastic Neighbor Embedding.

Data: data set $\mathcal{X} = \{x_1, x_2, \dots, x_n\}$,

cost function parameters: perplexity $Perp$,

optimization parameters: number of iterations T , learning rate η , momentum $\alpha(t)$.

Result: low-dimensional data representation $\mathcal{Y}^{(T)} = \{y_1, y_2, \dots, y_n\}$.

begin

compute pairwise affinities $p_{j|i}$ with perplexity $Perp$ (using Equation 1)

set $p_{ij} = \frac{p_{j|i} + p_{i|j}}{2n}$

sample initial solution $\mathcal{Y}^{(0)} = \{y_1, y_2, \dots, y_n\}$ from $\mathcal{N}(0, 10^{-4}I)$

for $t=1$ **to** T **do**

compute low-dimensional affinities q_{ij} (using Equation 4)

compute gradient $\frac{\delta C}{\delta \mathcal{Y}}$ (using Equation 5)

set $\mathcal{Y}^{(t)} = \mathcal{Y}^{(t-1)} + \eta \frac{\delta C}{\delta \mathcal{Y}} + \alpha(t) (\mathcal{Y}^{(t-1)} - \mathcal{Y}^{(t-2)})$


end

end


A selection from the 64-dimensional digits dataset

0	1	2	3	4	5	0	1	2	3	4	5	0	5						
5	5	0	4	1	3	5	1	0	0	2	2	0	1	2	3	3	3		
4	4	1	5	0	5	4	2	0	0	1	3	2	1	4	3	1	3	4	
3	4	1	4	0	5	3	1	5	4	9	2	2	2	5	5	4	0	0	1
2	3	4	5	0	1	2	3	4	5	0	1	2	3	4	5	0	5	5	5
0	4	1	3	5	1	0	0	2	2	0	1	2	3	3	3	4	4	4	4
1	5	0	5	2	2	0	0	1	3	2	1	3	1	3	4	3	1	4	4
0	5	7	4	5	4	4	1	2	2	5	5	4	0	0	1	2	3	4	4
5	0	1	2	3	4	5	0	4	2	3	4	5	0	5	5	5	0	4	1
3	5	4	0	0	2	2	0	4	2	3	3	3	4	4	1	5	0	5	0
5	2	2	0	0	1	3	2	4	3	1	3	1	4	3	1	4	0	5	5
3	1	5	4	2	2	2	5	5	4	0	3	0	1	2	3	4	5	0	5
0	1	2	3	4	5	0	1	2	3	4	5	0	5	5	5	0	4	1	3
5	1	0	0	1	2	2	0	1	2	3	3	3	4	4	1	5	0	5	0
1	2	0	0	1	3	2	4	3	1	3	1	4	3	1	4	0	5	3	3
1	5	4	4	2	2	2	5	5	4	0	0	1	2	3	4	5	0	1	3
2	3	4	5	0	1	2	3	4	5	0	5	5	5	0	4	1	3	5	4
0	0	1	2	2	0	1	2	3	3	3	4	4	1	5	0	5	2	2	2
0	0	1	3	2	1	4	3	1	3	1	4	3	1	4	0	5	3	1	5
4	4	2	2	1	5	5	4	4	0	0	1	2	3	4	5	0	1	3	5
0	0	1	2	2	0	1	2	3	3	3	4	4	1	5	0	5	2	2	2
0	0	1	3	2	1	4	3	1	3	1	4	3	1	4	0	5	3	1	5
4	4	2	2	1	5	5	4	4	0	0	1	2	3	4	5	0	1	3	5


Principal Components projection of the digits (time 0.03s)




MDS embedding of the digits (time 3.64s)




Isomap projection of the digits (time 1.44s)



Linear Discriminant projection of the digits (time 0.01s)




t-SNE embedding of the digits (time 15.92s)







(a) Visualization by t-SNE.



(b) Visualization by Sammon mapping.



(c) Visualization by Isomap.



(d) Visualization by LLE.






Figure 5: Visualizations of the COIL-20 data set.



ImageNet

http://cs.stanford.edu/people/karpathy/cnnembed/cnn_embed_4k.jpg



NIPS coauthorship

http://lvdmaaten.github.io/tsne/examples/nips_tsne.jpg